

Data-Semantics-Aware Recommendation of Diverse Pivot Tables

Abstract

Data summarization is essential to discover insights from large datasets. In spreadsheets, *pivot tables* offer a convenient way to summarize tabular data by computing aggregates over some attributes, grouped by others. However, identifying attribute combinations that will result in *useful* pivot tables remains a challenge, especially for high-dimensional datasets. We formalize the problem of automatically recommending *insightful* and *interpretable* pivot tables, eliminating the tedious manual process. A crucial aspect of recommending a *set* of pivot tables is to *diversify* them. Traditional works inadequately address the table-diversification problem, which leads us to consider the problem of *pivot table diversification*.

We present SAGE, a data-semantics-aware system for recommending *k*-budgeted diverse pivot tables, overcoming the shortcomings of prior work for top-*k* recommendations that cause redundancy. SAGE ensures that each pivot table is *insightful*, *interpretable*, and *adaptive* to the user’s actions and preferences, while also guaranteeing that the set of pivot tables are different from each other, offering a *diverse* recommendation. We make two key technical contributions: (1) a *data-semantics-aware model* to measure the utility of a single pivot table and the diversity of a set of pivot tables, and (2) a *scalable greedy algorithm* that can efficiently select a set of diverse pivot tables of high utility, by leveraging data semantics to significantly reduce the combinatorial search space. Our extensive experiments on three real-world datasets show that SAGE outperforms alternative approaches, and efficiently scales to accommodate high-dimensional datasets. Additionally, we present several case studies to highlight SAGE’s qualitative effectiveness over commercial software and Large Language Models (LLMs).

1 Introduction

Data is at the heart of data-driven decision making. We rely on data—specifically, trends observed in the data—to obtain *insights* [21] that help us make informed decisions. However, due to limitations in human comprehensibility, data must be *summarized* [27, 35, 59] to enable humans observe trends and discover insights from the summaries—either directly or via visualizations [62] over the summaries. One of the most common techniques to summarize data is *aggregation*. Simple aggregations involve functions (e.g., SUM) to aggregate all rows. More nuanced aggregations involve multiple groupings of the entities (e.g., GROUP BY GENDER, MARITAL_STATUS) and then aggregating each group separately.

While SQL provides functionalities for any custom aggregation query, it is not suitable for novices, and has various interface-related limitations. Thanks to the ubiquity of spreadsheet software—such as Microsoft Excel [44], Google Sheets [28], Apple Numbers [33], etc.—a substantial portion of businesses (around 60% [20]) and about 2 billion people [47] use spreadsheets for data management and analysis. These users rely on *pivot tables*, a summary of tabular data that computes aggregates over a few data attributes, grouped by other data attributes. Most commercial spreadsheets include a built-in and user-friendly mechanism to construct pivot tables. Spreadsheet pivot tables are particularly suitable for novices, where

ID	Gender	Age	Experience	Degree	Department	Salary
1	Male	48	3	PhD	IT	\$50,000
2	Female	32	1	MS	Sales	\$20,000
3	Male	45	12	PhD	HR	\$100,000

Figure 1: A sample table from an employee compensation dataset.

they can rearrange, group, and aggregate data using intuitive interfaces such as drag-and-drop. Unlike SQL aggregates, spreadsheet pivot tables offer dynamic user interactions, allowing interactive exploration such as drilling down, filtering, sorting, etc.

In an exploratory setting where the goal is to discover interesting data trends, a key challenge in constructing insightful pivot tables lies in selecting the right *parameters*, i.e., determining which attributes to use for groupings and aggregations. This task becomes even harder when users lack domain knowledge, face missing or cryptic attribute names, or work with high-dimensional data. In such cases, users must manually explore a vast space of parameter combinations through a tedious trial-and-error process. This involves experimenting with various combinations of (1) grouping attributes, (2) aggregation attributes, and (3) aggregation functions, then manually assessing the insightfulness of the resulting pivot tables. We illustrate this challenge with the following example.

EXAMPLE 1.1. *Sasha is investigating potential factors affecting salary across various groups in an employee compensation dataset over 21 attributes including ID, Gender, Age, Experience, Degree, Department, Salary, etc. (Figure 1). With an aim to discover salary discrepancies across various group combinations, she starts with the pivot table shown in Figure 2: she puts Gender in the row-groups (A) and Department in the column-groups (B); and chooses Salary as an aggregate attribute (C) and Average as the aggregate function (D).*

Sasha is interested in Salary discrepancies, so her choice for (C) is fixed. However, she still needs to explore various combinations for (A), (B), and (D). Sasha wishes to put demographic attributes (e.g., Gender, Race, Age, Marital Status, etc.) in the row-groups, as any discrepancy across different rows will indicate discrimination, and all other attributes in the column-groups. For this particular dataset, there are 5 demographic attributes and 15 non-demographic attributes. Sasha decides to explore 4 aggregation functions: MAX, MIN, AVERAGE, and SUM, as all of them can expose different types of discrepancies.

This leaves her $5 \times 15 \times 4 = 300$ possible combinations,¹ and she must carefully inspect each pivot table to identify salary discrepancies, by manually contrasting the pivot table cells. Assuming each pivot table has 10 cells on average and it takes about 2 minutes to examine each pivot table, Sasha needs $300 \times 2 = 600$ minutes (10 hours)!

Recommending Pivot Tables. Example 1.1 highlights the need for a smart *recommendation* system that can automatically suggest the “best” pivot tables, relieving Sasha from tedious, error-prone work. While existing spreadsheet software, such as Microsoft Excel and Google Sheets, are equipped with features for automatic pivot table recommendation, they have several shortcomings. We proceed to highlight some of their key limitations through Example 1.2.

¹Sasha chose only one option for each parameter. Multiple options (e.g., Gender and Race for row-groups) will further increase the search space of possible pivot tables.




Figure 2: A pivot table requires 4 parameters: (A) row-groups, (B) column-groups, (C) aggregate attributes, and (D) aggregate functions for each aggregate attribute. Users can choose multiple values for each parameter. In Example 1.1, Sasha has fixed (C), but still needs to explore (A) (5 options), (B) (15 options), and (D) (4 options).

EXAMPLE 1.2. Frustrated by the manual exploration, Sasha gives pivot table recommendation features of commercial spreadsheet software a try (Figure 3). Google Sheets provides three recommendations.² However, the recommended pivot tables often include too many aggregated attributes beyond what she wanted (Salary), resulting in convoluted and large pivot tables. Sasha also observes that most recommendations are redundant—they default to the groupings by Gender or Department—and lack diversity, causing her to miss out on insights involving other data attributes.

While Microsoft Excel provides nine recommendations,³ it utilizes only 10 out of 21 possible attributes. Additionally, it suggests meaningless aggregations such as Sum(Employed Year) and Sum(Age), revealing its shortcoming in grasping the data semantics. Furthermore, for both MS Excel and Google Sheets, Sasha failed to specify Salary as her intended aggregate attribute, restricting her ability to steer the recommendations towards her specific needs.

Lastly, Sasha asks ChatGPT for three “insightful” and “diverse” pivot tables, focusing on average Salary. While the recommendations initially seemed reasonable, she quickly realizes that the cell values of the pivot tables are entirely hallucinated,⁴ revealing that ChatGPT ignored the actual data and failed to validate the insightfulness of the pivot tables. Sasha concludes that LLMs are ill-suited for this task, since they generate outputs heuristically, not by explicitly enumerating and scoring all options for pivot tables.

Example 1.2 highlights several key limitations of existing tools for automatic recommendation of pivot tables. First, they do not cater to the user needs for a *focused* and *adaptive* recommendation of pivot tables. Second, they focus on top-k recommendations [19, 65, 66] and do not consider *diversification* [23], which may cause the users to miss certain data insights. Finally, existing approaches do not fully leverage the data and its semantics to ensure that the suggested pivot tables are *useful*, i.e., *insightful* and *interpretable*. We propose SAGE, a data-semantics-aware system for recommending k-budgeted diverse pivot tables, which overcomes the shortcomings of the existing approaches. We summarize the limitations of currently available tools and research work in Figure 4 to contrast them against SAGE, and defer a detailed discussion to Section 7.

²Test conducted in February 2025.

³Tested on Microsoft Excel (Windows) version 2501, February 2025.

⁴Tested on ChatGPT (GPT-4o), February 2025.

Tool	Recommended Pivot Tables
Google Sheets [28]	(1) Average of Age, Years of Experience, Annual Bonus, Overtime Hours, Sick Days, Training Hours, Satisfaction Score, #Projects, #Promotions by Gender
	(2) Average of Performance Rating for each Gender by Department
	(3) Average of Age, Years of Experience, Performance Rating, Salary, Annual Bonus, Overtime Hours, Sick Days, Training Hours, Satisfaction Score, #Projects, #Promotions by Gender
Microsoft Excel [44]	(1) Count of ID by Degree
	(2) Count of ID by Department
	(3) Sum of Employed Year by Department
	(4) Sum of #Promotions by Employed Year and Degree
	(5) Sum of Age, Children, Performance Rating by Degree
	(6) Sum of Children, Performance Rating, Salary by Degree
	(7) Sum of Children, Performance Rating, Salary by Department
	(8) Sum of Salary by Employed Year and Gender
	(9) Sum of Employed Year by Gender and Degree
ChatGPT [50]	(1) Average Salary by Years of Experience and Training Hours
	(2) Average Salary by Department and Children
	(3) Average Salary by Age and Satisfaction Score

Figure 3: Google Sheets recommendations are redundant and convoluted; Microsoft Excel includes meaningless recommendations such as SUM(Age); while ChatGPT recommendations look reasonable, they are data-content-unaware as the pivot table values are hallucinated.

Problem. The problem we study in this paper is recommending a diverse set of pivot tables, under a size constraint, while ensuring that each recommended pivot table is *useful*, meaning it is *insightful* and *interpretable*. Furthermore, we want to achieve two usability goals during recommendation: (1) *adaptivity*, which takes into consideration already explored pivot tables by the users, and (2) *customizability*, which enables the users to guide the recommendation process by specifying certain data attributes to prioritize.

Challenges. We now highlight three key challenges that are associated with the problem:

Challenge 1: semantic modeling of pivot table utility. A useful pivot table must be *insightful*, to inform the users interesting data trends, and *interpretable*, by not overwhelming the users with too much information. Traditional approaches [17, 21, 30, 62] ignore the *data semantics*; they use simple statistical measures to model insightfulness, which often fails to ensure interpretability and interestingness. For instance, high variance among artists across 100 hobby-related attributes might indicate a statistically significant trend, but is not very insightful (since it is expected that artists will have various hobbies). Furthermore, semantically modeling insightfulness and interpretability of a pivot table in a *multi-group setting* (e.g., group by Gender, Department, Degree) is non-trivial and is not addressed in prior work. In summary, how to model insightfulness and interpretability of a pivot table while remaining aware of the data semantics is a key challenge.

Challenge 2: modeling table diversity. Beyond recommending insightful and interpretable pivot tables, our goal is to also *diversify* the set of pivot tables. To the best of our knowledge, the notion of diversity in the context of pivot tables is not defined in prior work. Existing diversification approaches [11, 22–24] do not trivially extend for “table diversification”, where the items under consideration are entire tables rather than individual tuples. Prior works for recommending insightful data summaries [65] or visualizations [62] do not consider diversity. For pivot table diversification, the key challenge is to develop an appropriate distance metric to model both the syntactic (e.g., attribute coverage) and semantic (e.g., provided insights) distances between a pair of pivot tables.

		Commercial software	Research work	LLMs	This work				
		Microsoft Excel [44]							
		Google Sheets [28]							
		PowerBI [21, 45]							
		Tableau [57]							
		DAISY [65]	AutoSuggest [66]						
				ChatGPT [50]					
				Llama3-instruct [42]					
				TableGPT [60]					
					SAGE				
Desirable Properties	Budgeted recommendations	○	○	○	○	○	○	●	
	Guarantees syntactic validity	●	○	○	○	○	○	○	●
	Guarantees semantic validity	○	○	○	○	○	○	○	○
	Ensures interpretability	●	○	○	○	○	○	○	○
	Adaptive to user actions	○	○	○	○	○	○	○	○
	Allows user specifications	○	○	○	○	○	○	○	○
	Ensures diversity	○	○	○	○	○	○	○	○
	Attribute-name semantics aware	●	○	○	○	○	○	○	○
	Attribute-order insensitive	○	○	○	○	○	○	○	○
	Data-semantics aware	●	○	○	○	○	○	○	○
	No additional requirements	○	○	○	○	○	○	○	○
	Low-cost	●	○	○	○	○	○	○	○
Open-source	○	○	○	○	○	○	○	○	

Figure 4: SAGE satisfies all desirable properties. While PowerBI, Tableau, DAISY, and AutoSuggest do not directly/always recommend pivot tables, we include them due to their capability to recommend data summaries. Code/software for DAISY and AutoSuggest is unavailable, thus we obtain their properties from the papers, and mark certain things as unknown. LLMs are not designed to directly recommend pivot tables, but they can be prompted to do so.

Challenge 3: developing an efficient system. Our goal is to recommend highly insightful and interpretable pivot tables, while ensuring diversity among them. Unlike insightfulness and interpretability, which can be measured for each pivot table in isolation, diversity requires considering a *set* of pivot tables. This leads to an NP-hard combinatorial search problem. Furthermore, evaluating insightfulness of a pivot table requires its materialization, which adds to the computational complexity. While greedy approaches with approximation guarantees [11, 13, 23, 48, 52] can alleviate the problem of combinatorial search, the requirement of materializing candidate pivot tables remain. Even with approximation algorithms [3, 15, 32, 64] for efficient materialization, without aggressive pruning before materialization, far too many candidates become the bottleneck. Therefore, a key challenge here is to develop mechanisms that can leverage semantic understanding of the data to prune unpromising pivot tables and avoid unnecessary materialization. Another challenge is to discover effective techniques that can “push down” [67] components of the diversity requirements to the search process to further prevent unnecessary pivot table materialization.

Contributions. Our main contribution is development of a novel system SAGE, for recommendation of a diverse set of useful pivot tables under a budget (size) constraint. Below, we provide the key contributions we make in this paper:

- We motivate and *formalize the problem* of budgeted recommendation of a diverse set of useful pivot tables, model it as a *constrained optimization problem*, and establish its *desiderata* (Section 2).
- We provide a formal model to measure the *utility* of a pivot table in terms of insightfulness and interpretability. Unlike prior work, our utility model leverages data semantics (Section 3).
- We establish the notion of *pivot-table diversification*, a key component of the problem we study in this paper. Our contribution lies

in the formulation of a suitable *distance metric*—which considers both structural and semantic properties of pivot tables—and its application to diversifying a set of pivot tables (Section 4).

- To ensure SAGE’s efficiency and practicality, we must tackle the NP-hardness of the problem. To reduce the search space, we introduce *aggressive semantic pruning*. To expedite the recommendation process, we leverage offline computation and “push down” diversity requirements to the search process (Section 5).
- We present a thorough empirical analysis over 3 real-world datasets and present case studies to qualitatively contrast SAGE against commercial software and LLMs. We show that SAGE can recommend diverse and high-utility pivot tables, outperforming prior approaches and other variants. We also find SAGE scalable and efficient in practice (Section 6).

2 Recommending a Diverse set of Pivot Tables

In this section, we motivate the need for diversity and adaptivity during recommending pivot tables (Section 2.1). Then we develop the desiderata for the problem (Section 2.2) and formalize it (Section 2.3). Finally, we provide an overview of SAGE (Section 2.4).

2.1 The Need for Diversity and Adaptivity

A key limitation of top-k recommendation is that it may provide redundant information, causing the users to miss out on relatively less useful, but complementary data insights. Such lack of *diversity* may even mislead the users to believe in partial insights that are “half-true”. Another shortcoming of existing approaches is that they are not *adaptive* to user actions, i.e., when the user acknowledges a recommendation, it should be excluded from the subsequent iterations. However, commercial pivot table recommendation features are not adaptive (Figure 4). We proceed to provide an example to highlight the need for *diverse* recommendation to help the users get a broader picture of the dataset, and *adaptive* [39, 58] recommendation to enable the user guide the recommendation process.

EXAMPLE 2.1. Recall from Example 1.1 that Sasha is interested in salary discrepancies. She initially finds two pivot tables (Figure 5 (a) & (b)) suggesting gender-based pay gap. However, a deeper pattern emerges when she expands her analysis using other aggregate functions (e.g., COUNT) and discovers the pivot tables shown in Figure 5 (c) & (d), which provide her an additional context that the discrepancy stems from the hiring process: employee counts are uneven across degrees and departments. Sasha also notes that IT employees earn more than those in Sales, and PhDs earn more than others. This indicates degree- and department-based discrepancies, which are expected and acceptable. Sasha concludes that males earn more on average not due to gender bias, but largely because more male PhDs work at IT.⁵

Furthermore, in an incremental setting where Sasha iteratively requests for recommendations of a few pivot tables at a time, she expects the system to adapt to her actions. For instance, after she accepts or rejects the suggestions of Figure 5 (a) and (b), the system should avoid recommending redundant pivot tables that reiterate the same concept (gender-based salary gap) across other aspects (e.g., marital status).

⁵This phenomenon is commonly known as Simpson’s paradox [9, 10]. While our focus is not to explicitly expose Simpson’s paradox, we show this as a motivating use-case to highlight the need for diversification in pivot table recommendations.

Gender	Degree			Gender	Department	
	BS	MS	PhD		IT	Sales
Male	200K	300K	1000K	Male	1000K	500K
Female	100K	200K	300K	Female	400K	200K

(a) Avg. Salary by Gender and Degree

Gender	Degree, Department			Gender	Department	
	BS	MS	PhD		IT	Sales
Male	4	1	8	Male	1000K	500K
Female	2	8	2	Female	400K	200K

(b) Avg. Salary by Gender and Dept.

Gender	Degree, Department			Gender	Department	
	BS	MS	PhD		IT	Sales
Male	4	1	8	Male	1000K	500K
Female	2	8	2	Female	400K	200K

(c) Count ID by Gender, Degree, and Dept.

Degree	Department	
	IT	Sales
BS	200K	100K
MS	300K	200K
PhD	900K	400K

(d) Avg. Salary by Degree and Dept.

Figure 5: Four pivot tables over the dataset of Figure 1. While (a) and (b) indicate gender-based salary gap, (c) and (d) add additional context.

2.2 Desiderata

We now provide five key desiderata for an ideal system for recommending a set of pivot tables:

- D1. Each recommended pivot table must provide *insightful* [17, 62] and *semantically interesting* information. For instance, a significant gap in average salary across genders provides insight into gender-based pay gap. Furthermore, “average salary” is semantically interesting, where “sum of zip code” is not.
- D2. Each pivot table must be *interpretable*, ensuring ease of comprehension by humans. For instance, a concise table with 10 cells is more interpretable than one with 1000 cells.
- D3. While insightfulness and interpretability model the goodness of a single pivot table, a desirable property for a set of pivot tables is *diversity*. Thus, the recommended set of pivot tables must minimize redundancy, covering various data aspects.
- D4. The system for pivot table recommendation must allow (I) *customizability*—allowing users to specify the desired size of the recommendation set, degree of diversity, data scope, etc.—and (II) *adaptiveness* to user actions.
- D5. Finally, the system must be *efficient* and *scalable*—to ensure handling large and high-dimensional data effectively—and *accessible*—in terms of cost and availability.

2.3 Problem Formulation

We now formalize our problem for a single-relation database instance (dataset) D by defining a pivot table.

Definition 2.1 (Pivot Table). Given a dataset D over a set of attributes A and the domain of aggregate functions $\mathcal{F} = \{\text{COUNT}, \text{SUM}, \text{AVG}, \text{MIN}, \text{MAX}\}$, a pivot table $T(F(V), G)$ takes the following form:

SELECT $F(V)$ FROM D GROUP BY G

where, $G = \{G_1, G_2, \dots\} \subseteq A$ is a subset of attributes for grouping; $V = \{V_1, V_2, \dots\} \subseteq A$ is a subset of attributes for computing aggregates over; $G \cap V = \emptyset$ ensures that no attribute is used for both grouping and aggregation; $F = \{F_1, F_2, \dots\}$ is a set of aggregate functions where $F_i \in \mathcal{F}$ and $|F| = |V|$; and with slight abuse of notation, $F(V)$ denotes $F_1(V_1), F_2(V_2), F_3(V_3) \dots$

Tabular representation of a pivot table. For $T(F(V), G)$, with $|G| \geq 2$, we fix as row-groups and column-groups (Figure 2) two non-empty sets $R, C \subset G$, respectively, where $R \cup C = G$, $R \cap C = \emptyset$.

EXAMPLE 2.2. Figure 5(c) represents a possible tabular representation for the pivot table: SELECT COUNT(ID) FROM D GROUP BY Gender, Degree, Department. Here, $R = \{\text{Gender}\}$ and $C = \{\text{Degree, Dept.}\}$

Pivot table canonicalization. The above mechanism allows structurally different tabular representations of semantically equivalent pivot tables. However, our focus is on the semantics of a pivot table and not the specific orientation of its tabular representation. Thus,

we canonicalize a pivot table $T(F(V), G)$ by lexicographically sorting $F(V)$ and G to obtain $F(V)_{\leq}$ and G_{\leq} , respectively, and derive the canonical pivot table $T(F(V)_{\leq}, G_{\leq})$. Furthermore, we obtain a *canonical tabular representation* of a pivot table $T(F(V), G)$ by assigning to the row-groups (R_{\leq}) the first $\lceil \frac{|G|}{2} \rceil$ elements of G_{\leq} and the remaining elements to the column-groups (C_{\leq}). This process ensures that pivot tables remain *organization-invariant* (e.g., transpose-invariant), i.e., $T(F(V), G)$ and all its variants derived from different permutations of $F(V)$ and G result in an identical canonical tabular representation.

EXAMPLE 2.3. The canonical pivot table for Figure 5(c) is: SELECT COUNT(ID) FROM D GROUP BY Degree, Department, Gender. The canonical tabular representation is obtained by setting $R_{\leq} = \langle \text{Degree, Dept.} \rangle$ and $C_{\leq} = \langle \text{Gender} \rangle$, which is simply the transpose of Figure 5(c).

Based on the desiderata of Section 2.2, we set our goal to find a bounded sized (D4-I) set of pivot tables such that the overall utility (D1 & D2) of the pivot tables are maximized while the set of pivot tables meet the minimum diversity requirement (D3).

PROBLEM 2.1 (RECOMMENDING A SET OF PIVOT TABLES). Given (i) a set of possible pivot tables \mathcal{T}_A over a dataset D with attributes A , (ii) a function $Utility : \mathcal{T}_A \mapsto [0, 1]$ that returns the utility of a pivot table $T \in \mathcal{T}_A$, (iii) a function $Diversity : 2^{\mathcal{T}_A} \mapsto [0, 1]$ that returns the diversity of a set of pivot tables $T \subseteq \mathcal{T}_A$, (iv) a budget $k \in \mathbb{N}^+$, and (v) a threshold $\theta \in [0, 1]$, find a set of pivot tables $T^* \subseteq \mathcal{T}_A$ s.t:

$$\begin{aligned}
 &(\text{objective}) && T^* = \arg \max_{T \subseteq \mathcal{T}_A} \sum_{T \in T} Utility(T), \\
 &(\text{size constraint}) && |T^*| \leq k, \text{ and} \\
 &(\text{diversity constraint}) && Diversity(T^*) \geq \theta
 \end{aligned}$$

Problem 2.1 balances utility and diversity by maximizing utility while putting a constraint on diversity. Other variants of this problem are possible such as maximizing a linear combination of the objective and the diversity constraint. More details are in Section 5.

Adaptive Recommendation of a set of Pivot Tables. In the adaptive version (D4-II), we discard the already explored pivot tables by the user T_u from the set \mathcal{T}_A to obtain $\mathcal{T}_A - T_u$. When the user highlights a data scope (D4-I) by specifying a subset of attributes $A_u \subseteq A$ they want to focus on, we set the possible pivot tables to \mathcal{T}_{A_u} .

Considerations. Multiple aggregates within a pivot table is essentially equivalent to concatenating the corresponding single-aggregate pivot tables, i.e.,

$$T(F(V), G) \equiv \bigcup_{F, V \in F, V} T(F(V), G)$$

Therefore, for simplicity and to promote interpretability, we limit each pivot table to have exactly one aggregate. We use F and V to denote the aggregation function and attribute, respectively. We summarize the notations used in the rest of this paper in Figure 6.

Symbol	Description
$D, A, D[A]$	Database, attributes, possible unique value combinations
G, R, C	Grouping attributes, row-groups, column-groups; $R \cup C = G$
F, V	Aggregation function and attribute
$T(F(V), G)$ or T	A pivot table for the query <code>SELECT F(V) FROM D GROUP BY G</code>
T_{r_i}/T^{c_j}	The row/column of T with row/column header r_i/c_j
$T_{r_i}^{c_j}$	The pivot table cell with row header r_i and column header c_j
n, m	The cardinality of $D[R]$, $D[C]$

Figure 6: Table of notations. We use bold letters to denote sets.

A Note on Generalizability. While our work focuses on recommending pivot tables in spreadsheet environments, the techniques can be generalized to recommend aggregate queries in relational databases. Pivot tables can be represented by SQL aggregate queries involving Group-by, allowing SAGE’s adaptation in RDBMS.

2.4 SAGE Overview

To model utility of a single pivot table, SAGE utilizes two properties, *insightfulness* and *interpretability*, covering both syntactic and semantic aspects (Section 3). SAGE models diversity based on the *semantic distance* between a pair of pivot tables (Section 4). SAGE performs some offline precomputation, based on which, it applies an online greedy algorithm to generate a set of diverse and useful pivot tables as a solution for Problem 2.1 (Section 5).

3 Utility of a Pivot Table

In this section, we describe how we quantify the goodness or *utility* of a single pivot table. Based on the desiderata of Section 2.2, a pivot table has high utility if it offers *insights* (D1) while being easily *interpretable* by humans (D2). Thus, we use insightfulness (Section 3.1) and interpretability (Section 3.2) as the two primary building blocks to model the utility of a pivot table.

3.1 Insightfulness

Intuitively, an insightful pivot table must involve attributes that are *significant*, i.e., inherently interesting and relevant (§ 3.1.1). Furthermore, it should satisfy at least one of the following criteria: (1) provide high *informativeness* (§ 3.1.2), (2) highlight meaningful *trends* (§ 3.1.3), or (3) reveal *surprising* [14, 34, 59] findings (§ 3.1.4). We build on prior works [14, 17, 21, 30, 34, 62] that model insightfulness based on only statistical properties, but significantly extend it by taking a *semantics-aware* approach, enabled by LLMs [42].

3.1.1 Attribute significance. Typically, not all data attributes are of interest by human users. For instance, grouping data by Name is typically much less insightful than by Gender. However, semantic understanding is required to figure out attribute significance. To this end, we consult an LLM [42] to determine the significance for an attribute A . When attribute name is missing or semantically meaningless (e.g., “Column 1”), we first query an LLM to suggest appropriate names for attributes by providing it with a small sample of the data. LLM’s semantic-reasoning capability allows us to achieve this without any domain-specific pre-configuration. To avoid noise, we employ multiple paraphrased prompts while querying the LLM. In this work, we condition the LLM to return a simple binary answer (yes \rightarrow 1/no \rightarrow 0). However, this component can be replaced by a domain-aware model that can return the likelihood of an attribute being significant for a specific context. We compute attribute significance of a pivot table T , $S_{\text{sig}} : \mathcal{T}_A \mapsto [0, 1]$

as follows:

$$S_{\text{sig}}(T) = \prod_{A \in \{V\} \cup G} \text{Significance}(A) \quad (1)$$

Here, $\text{Significance} : A \mapsto [0, 1]$ denotes the probability that an attribute $A \in A$ is a significant attribute w.r.t human interest.

EXAMPLE 3.1. For Figure 5(d), Degree, Department, and Salary, all are significant attributes. Thus, $S_{\text{sig}}(T) = 1 \times 1 \times 1 = 1$.

3.1.2 Informativeness. A statistical way to measure informativeness within data is to measure spread of the values. Intuitively, if values in a pivot table deviate from each other significantly, the “entropy” is high, and so is the informativeness. In this work, we use *deviation* across different groups to model informativeness. Unlike prior works [17, 30, 62] that only consider a two-group setting (Male vs Female), we consider a multi-group setting.

Given a database D and a pivot table T with row-groups R and column-groups C , let $D[R]$ and $D[C]$ be the set of row and column headers for T , respectively. E.g., for Figure 5(c), the row headers are {Male, Female} and the column headers are {(BS, IT), (BS, Sales), (MS, IT), (MS, Sales), (PhD, IT), (PhD, Sales)}. We use T_{r_i} (T^{c_i}) to denote the row (column) of T with row (column) header r_i (c_i). We use n and m to denote the number of rows $|D[R]|$ and columns $|D[C]|$ in T , respectively. We compute the row-wise and column-wise informativeness scores $S_{\text{inf}}^{\text{row}}$ and $S_{\text{inf}}^{\text{col}}$ as follows:

$$S_{\text{inf}}^{\text{row}}(T) = \frac{1}{\binom{n}{2}} \sum_{r_i, r_j \in D[R] \text{ s.t. } i < j} \frac{\|T_{r_i} - T_{r_j}\|_2}{\gamma \cdot m}$$

$$S_{\text{inf}}^{\text{col}}(T) = \frac{1}{\binom{m}{2}} \sum_{c_i, c_j \in D[C] \text{ s.t. } i < j} \frac{\|T^{c_i} - T^{c_j}\|_2}{\gamma \cdot n}$$

Here, γ is a normalization parameter, set to $\max(T) - \min(T)$, ensuring that $S_{\text{inf}}^{\text{row}}$ and $S_{\text{inf}}^{\text{col}}$ are bounded between 0 and 1. Also note that while we use Euclidean distance (L_2 distance), any other distance function such as L_1 distance can be used here. We compute the informativeness score $S_{\text{inf}} : \mathcal{T}_A \mapsto [0, 1]$ by taking the maximum of the row-wise and column-wise informativeness scores:

$$S_{\text{inf}}(T) = \max(S_{\text{inf}}^{\text{row}}(T), S_{\text{inf}}^{\text{col}}(T)) \quad (2)$$

EXAMPLE 3.2. We first compute the pairwise distances along the rows of Figure 5(d): $\|T_{BS} - T_{MS}\|_2 = 141.4K$, $\|T_{BS} - T_{PhD}\|_2 = 761.6K$, and $\|T_{MS} - T_{PhD}\|_2 = 632.5K$. We normalize using $\gamma = 900K - 100K = 800K$ and $m=2$, resulting in normalized distances of [0.088, 0.476, 0.395]. Taking an average gives us $S_{\text{inf}}^{\text{row}}(T) = 0.32$. We similarly compute $S_{\text{inf}}^{\text{col}}(T) = 0.22$ and obtain $S_{\text{inf}}(T) = \max(0.32, 0.22) = 0.32$.

3.1.3 Trend. Trends observed in a pivot table provide insights. However, the degree of insightfulness hinges on two key factors: the *magnitude* of the trend metric and how *atypical* or rare it is. For instance, a positive correlation between income and years of service is generally expected—employees with longer tenures typically earn more. In contrast, a trend showing that new hires earn more on average than long-serving employees contradicts this expectation and thus is particularly insightful.

We use two metrics to quantify the magnitude of a trend: *correlation* and *ratio*. Furthermore, to assess the degree of a trend’s rarity, we query an LLM, which is aware of a broader semantic context. Thus, our definition of the *trend score* for a pivot table combines (1) purely statistical insights, reflected in high correlation and

consistent ratio across pivot table values and (2) semantic insights, captured through the LLM’s assessment of the trend’s atypicality.

Correlation. We use $\rho_{i,j}$ to denote the Pearson correlation coefficient between the rows T_{r_i} and T_{r_j} . Since consulting LLMs is costly, we only consider significant correlations and require the magnitude to be at least τ_ρ , a customizable threshold parameter, with a default value of 50%. The indicator function $[[|\rho_{i,j}| \geq \tau_\rho]]$ denotes if the correlation between the rows T_{r_i} and T_{r_j} is significant. We compute the row-wise correlation-trend score $S_{\text{cor}}^{\text{row}} : \mathcal{T}_A \mapsto [0, 1]$ as follows:

$$S_{\text{cor}}^{\text{row}}(T) = \frac{1}{\binom{n}{2}} \sum_{r_i, r_j \in D[R] \text{ s.t. } i < j} |\rho_{i,j}| \cdot [[|\rho_{i,j}| \geq \tau_\rho]] \cdot \widetilde{Pr}_{\text{cor}}(r_i, r_j)$$

Here, $\widetilde{Pr}_{\text{cor}}(r_i, r_j)$ is the likelihood of *not* observing a high correlation between the groups r_i and r_j determined by an LLM. We prompt an LLM “In a five-point scale from *very likely* to *very unlikely*, how likely is it that the $\langle F(V) \rangle$ for $\langle r_i \rangle$ and $\langle r_j \rangle$ are \langle positively/negatively \rangle correlated across $\langle D[C] \rangle$?”, where each part within $\langle \rangle$ is replaced with actual values such as “Average Salary” for $\langle F(V) \rangle$. We then map the LLM-provided likelihood to a numerical value using the mappings: *very likely* \rightarrow 20%, *likely* \rightarrow 40%, *neutral* \rightarrow 60%, *unlikely* \rightarrow 80%, and *very unlikely* \rightarrow 100%. The intuition behind this *inverse* mapping is that the more unlikely it is to observe a trend, the more insightful it is. We compute $S_{\text{cor}}^{\text{col}}(T)$ similarly and set the correlation-trend score $S_{\text{cor}}(T) = \max(S_{\text{cor}}^{\text{row}}(T), S_{\text{cor}}^{\text{col}}(T))$.

EXAMPLE 3.3. In Figure 5(d), T_{BS} and T_{PhD} exhibit a positive correlation of 98%, which is likely (from LLM consultation), leading $Pr_{\text{cor}}(\text{BS}, \text{PhD})$ to be 40%. The correlation between T_{BS} and T_{MS} is 100% (very likely); and T_{MS} and T_{PhD} is 100% (likely). Since all correlations meet the threshold 50%, we compute $S_{\text{cor}}^{\text{row}}(T) = (0.98 \times 0.4 + 1.0 \times 0.2 + 1.0 \times 0.4) / 3 = 0.33$. The column-wise correlation-trend score $S_{\text{cor}}^{\text{col}}(T) = (0.98 \times 0.4) / 1 = 0.39$. Thus, $S_{\text{cor}}(T) = \max(0.33, 0.39) = 0.39$.

Ratio. Since correlation fails to capture the relative *magnitude*, we use *ratio trends* based on *persistent* ratios between two groups, e.g., T_{PhD} earning at least $5\times$ more than T_{BS} across all departments. Similar to correlation trends, LLMs inform us the rarity of ratio trends. We compute the row-wise ratio-trend score $S_{\text{ratio}}^{\text{row}} : \mathcal{T}_A \mapsto [0, 1]$ as follows:

$$S_{\text{ratio}}^{\text{row}}(T) = \frac{1}{\binom{n}{2}} \sum_{r_i, r_j \in D[R]} (1 - \frac{1}{\pi_{i,j}}) \cdot [[\pi_{i,j} \geq \tau_\pi]] \cdot \widetilde{Pr}_{\text{ratio}}(r_i, r_j)$$

Here, $\pi_{i,j}$ denotes the minimum element-wise ratio between T_{r_i} and T_{r_j} , i.e., the smallest factor by which any value in T_{r_i} exceeds its corresponding value in T_{r_j} . To reduce LLM consultation cost, we use a threshold τ_π and require $\pi_{i,j}$ to be at least τ_π before consulting an LLM. While we set $\tau_\pi = 2.0$, it is a customizable parameter (but must be ≥ 1). $\widetilde{Pr}_{\text{ratio}}(r_i, r_j)$ denotes the LLM-provided likelihood of *not* observing the ratio trend between T_{r_i} and T_{r_j} . Note that for any i and j , at most one of $\pi_{i,j}$ or $\pi_{j,i}$ can contribute to this score, hence we fix the scaling factor to $\binom{n}{2}$. We normalize the trend magnitude by subtracting the inverse of $\pi_{i,j}$ from 1, so that larger ratios yield higher scores. We compute the column-wise ratio-trend score $S_{\text{ratio}}^{\text{col}}(T)$ similarly and set the ratio-trend score $S_{\text{ratio}}(T) = \max(S_{\text{ratio}}^{\text{row}}(T), S_{\text{ratio}}^{\text{col}}(T))$.

EXAMPLE 3.4. In Figure 5(d), the minimum ratio between T_{MS} and T_{BS} is 1.5; between T_{PhD} and T_{BS} is 4.0; and between T_{PhD} and

T_{MS} is 2.0. After applying the threshold $\tau_\pi = 2.0$, the ratio trends for $(T_{\text{PhD}}, T_{\text{MS}})$ and $(T_{\text{PhD}}, T_{\text{BS}})$ are retained. The LLM returns the likelihoods: [Very Unlikely, Unlikely] \rightarrow [1.0, 0.8] for these two trends. This gives us the row-wise ratio-trend score $S_{\text{ratio}}^{\text{row}}(T) = (3/4 \times 1.0 + 1/2 \times 0.8) / 3 = 0.37$. For the column-wise ratio-trend score, no pair satisfies the threshold requirement and thus the score is 0.0. Therefore, the ratio-trend score $S_{\text{ratio}}(T) = \max(0.37, 0.0) = 0.37$.

Finally, we compute the trend score by taking the maximum of the correlation-trend and ratio-trend scores:

$$S_{\text{trend}}(T) = \max(S_{\text{cor}}(T), S_{\text{ratio}}(T)) \quad (3)$$

EXAMPLE 3.5. For the pivot table of Figure 5(d), we computed the correlation-trend score as 0.39 in Example 3.3 and the ratio-trend score as 0.37 in Example 3.4. Thus, $S_{\text{trend}}(T) = \max(0.39, 0.37) = 0.39$.

3.1.4 Surprise. Surprising values or *outliers* often indicate existence of insights. E.g., in Figure 5(d), $T_{\text{PhD}}^{\text{IT}}$ is exceptionally high (900K) in its column. While such outliers can be insightful, not all are. Some, like this one, are expected: IT is high-paying, and PhDs earn more. In contrast, an unusually high $T_{\text{BS}}^{\text{Sales}}$ would be surprising, and thus insightful. Beyond simply identifying outliers, we incorporate the unexpectedness of observing outliers using LLM’s semantic knowledge. We compute the row-wise surprise score $S_{\text{sur}}^{\text{row}} : \mathcal{T}_A \mapsto [0, 1]$ as follows:

$$S_{\text{sur}}^{\text{row}}(T) = \frac{1}{n} \sum_{r_i \in D[R]} \text{OutlierScore}(T_{r_i})$$

$$\text{where, } \text{OutlierScore}(T_{r_i}) = \begin{cases} 1 - \frac{\sum_{c \in O_{r_i}} \widetilde{Pr}_{\text{outlier}}(r_i, c, T_{r_i}^c)}{|O_{r_i}| + 1}, & \text{if } |O_{r_i}| > 0 \\ 0, & \text{otherwise} \end{cases}$$

$$O_{r_i} = \{c_j \in D[C] \text{ s.t. } |\pi_{r_i}^{c_j} - \mu(T_{r_i})| \geq \tau_O \cdot \sigma(T_{r_i})\}$$

O_{r_i} is the set of column headers for each outlier in T_{r_i} . E.g., $O_{\text{PhD}} = \{\text{IT}\}$, if 900K is an outlier for the row T_{PhD} . We ensure that the score increases with the number of outliers by subtracting the inverse of their count from 1. However, we believe that even a single outlier should contribute meaningfully to the score. To reflect this, we add 1 to the denominator so that having only one outlier results in a score multiplier of 0.5. The threshold τ_O is set to 4, since, assuming normal distribution, 99.99% of the population is expected to lie within 4 standard deviations from the average [25]⁶; and anything outside this range is an outlier. $\widetilde{Pr}_{\text{outlier}}(r_i, c, T_{r_i}^c)$ denotes the LLM-obtained likelihood of $T_{r_i}^c$ *not* being an outlier w.r.t T_{r_i} . We compute $S_{\text{sur}}^{\text{col}}(T)$ similarly and compute the surprise score:

$$S_{\text{sur}}(T) = \max(S_{\text{sur}}^{\text{row}}(T), S_{\text{sur}}^{\text{col}}(T)) \quad (4)$$

Computing Insightfulness. A pivot table is considered insightful if it exhibits any of the key characteristics: informativeness, trend, or surprise. Therefore, we take the maximum of the three scores— S_{inf} , S_{trend} , and S_{sur} —as computed in Equations 2, 3, and 4 to compute *Insightfulness* : $\mathcal{T}_A \mapsto [0, 1]$. This approach prioritizes the strongest signal among the three insight indicators. To prioritize pivot tables involving significant attributes, we scale this score by S_{sig} (Equation 1).

$$\text{Insightfulness}(T) = S_{\text{sig}}(T) \cdot \max(S_{\text{inf}}(T), S_{\text{trend}}(T), S_{\text{sur}}(T)) \quad (5)$$

⁶Data in pivot tables may not be normally distributed. For simplicity, we use this heuristic for outlier detection. Our contribution is not inventing outlier-detection methods and the user is free to replace this with more suitable methods.

EXAMPLE 3.6. For the pivot table of Figure 5(d), the attribute significance score $S_{\text{sig}}(T)=1$ (Example 3.1). The informativeness score $S_{\text{inf}}(T)=0.32$ (Example 3.2), the trend score $S_{\text{trend}}(T)=0.32$ (Example 3.5), and since there is no outlier in the pivot table, the surprise score $S_{\text{sur}}(T)=0.0$. Thus $\text{Insightfulness}(T)=1 \times \max(0.32, 0.39, 0.0)=0.39$.

Remark. Presence of outliers can inflate the normalization factor γ (Section 3.1.2), causing S_{inf} to shrink significantly due to the compression of the value range of non-outliers. However, such cases typically yield a high S_{sur} , complementing the low S_{inf} . Since the *Insightfulness* score is the maximum of the three components, a strong signal from any source suffices to indicate insightfulness.

3.2 Interpretability

While insightfulness is a key measure of a pivot table’s utility, it does not take into account the cognitive constraints of human users, who require *interpretability*. Consider the pivot table in Figure 7, which shows SUM(AGE) grouped by Degree, Employed Year, and Department. Its interpretability suffers due to: (i) high sparsity resulting from many value combinations yielding empty sets, such as no MS hire in IT in 2011, yielding *nulls* after aggregation (§ 3.2.1), (ii) semantically invalid aggregate SUM(AGE) (§ 3.2.2), and (iii) excessive columns from fine-grained yearly grouping, which compromises conciseness of the pivot table (§ 3.2.3). We proceed to describe three desirable interpretability properties of a pivot table.

3.2.1 Density. Each pivot table cell maps to a data subset under a specific value combination (e.g., MS hires in IT in 2011), so empty subsets are expected. When aggregated, these empty subsets produce *null* values. However, excessive *nulls* hinder interpretability [1], as humans struggle to draw insights from sparse tables. A common workaround is imputation with zeros. However it is misleading, as *nulls* denote missing data, whereas zeros may imply valid values (e.g., 0 for average temperature suggests an actual recording). This motivates a key interpretability criterion: high *density*. We compute the *density score* $S_{\text{den}} : \mathcal{T}_A \mapsto [0, 1]$ as follows:

$$S_{\text{den}}(T) = \frac{\sum_{(r_i, c_j) \in D[R] \times D[C]} \mathbb{I}[T_{r_i}^{c_j} \neq \text{null}]}{n \cdot m} \quad (6)$$

EXAMPLE 3.7. The pivot table of Figure 5(d) has 3 rows and 2 columns (total 6 cells) and no null values. Hence, $S_{\text{den}}(T) = \frac{6}{2 \times 3} = 1.0$.

3.2.2 Semantic validity. Row and column headers in a pivot table represent unique values of the grouping attributes in G . For interpretability, these headers must be semantically meaningful [8]. E.g., Degree with values {MS, BS, PhD} is interpretable, while a functionally equivalent Degree_ID with values {1, 2, 3} is not, due to the lack of direct semantic meaning [16]. Similarly, the aggregate function F must be semantically valid w.r.t V : AVG is semantically valid for AGE, but SUM is not [37]. Though intuitive for humans, such judgments require domain knowledge. Thus, we leverage an LLM to mimic human reasoning and assess aggregation semantics. We define the *semantic validity score* of $T(F(V), G)$ based on two criteria: (1) whether the data types of attributes in G are textual, and (2) the extent to which F is semantically valid w.r.t V .

$$S_{\text{sem}}(T) = \frac{|\{A \in G \text{ s.t. } \text{DataType}(A) \text{ is Text}\}|}{|G|} \cdot \text{Pr}_{\text{agg}}(F, V) \quad (7)$$

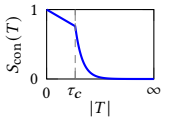
Degree	Employed Year					
	2011		2012		2024	
	IT	Sales	IT	Sales	IT	Sales
BS	428	304	251	null	...	256 192
MS	null	null	null	null	...	null null
PhD	603	650	441	450	...	null null

Figure 7: pivot table for the query SELECT SUM(AGE) GROUP BY EMPLOYED_YEAR, DEGREE with several interpretability issues.

We compute $\text{Pr}_{\text{agg}}(F, V)$ from an LLM-generated ranking of $F \in \mathcal{F}$ based on its semantic validity w.r.t V . We score the best function 1.0, the next 0.8, and so on, which ensures that $S_{\text{sem}}(T) \in [0, 1]$.

EXAMPLE 3.8. The pivot table of Figure 5(d), contains only textual headers. LLM responds to our prompt “Rank the functions COUNT, AVG, SUM, MIN, and MAX, based on their appropriateness for analyzing Salary” with {AVG, ...}. Thus, $S_{\text{sem}}(T) = \frac{2}{2} \times 1.0 = 1.0$.

3.2.3 Conciseness. While multiple grouping attributes may enhance a pivot table’s insightfulness, too many cells reduce its comprehensibility [51], and, thus, interpretability. To model this, we define conciseness score $S_{\text{con}} : \mathcal{T}_A \mapsto [0, 1]$ using a piecewise function [18]:

$$S_{\text{con}}(T) = \begin{cases} 1 - z|T|, & \text{if } |T| \leq \tau_c \\ (1 - z\tau_c)e^{-\lambda(|T| - \tau_c)}, & \text{if } |T| > \tau_c \end{cases} \quad (8)$$


Here, $|T|$ denotes the number of cells in T . This formula captures the intuition that interpretability declines gradually at first, but drops sharply once $|T|$ exceeds a threshold τ_c , set to 16. We apply a 3% linear decrease ($z = 0.03$) until $|T|$ exceeds τ_c , and an exponential decay at a rate of 50% ($\lambda = 0.5$) beyond that. This is grounded in cognitive load theory [7, 46], which states that performance declines sharply when cognitive demand exceeds working-memory capacity.

EXAMPLE 3.9. The pivot table of Figure 5(d) has 6 cells. Since $6 < 16$, we compute the linear part: $1 - 0.03 \times 6 = 0.82$. Thus, $S_{\text{con}}(T) = 0.82$.

Computing Interpretability. Unlike insightfulness, where a strong signal from any single type of insight is sufficient, interpretability demands that *all* criteria be met simultaneously. Therefore, we compute *Interpretability* : $\mathcal{T}_A \mapsto [0, 1]$ as the average of the three scores: S_{den} , S_{sem} , and S_{con} (Equations 6, 7, and 8):

$$\text{Interpretability}(T) = \frac{S_{\text{den}}(T) + S_{\text{sem}}(T) + S_{\text{con}}(T)}{3} \quad (9)$$

EXAMPLE 3.10. For Figure 5(d), we obtained values for $S_{\text{den}}(T)$, $S_{\text{sem}}(T)$, and $S_{\text{con}}(T)$ to be 1.0, 1.0, and 0.82, respectively, in the previous examples. Thus, $\text{Interpretability}(T) = (1.0 + 1.0 + 0.82)/3 = 0.94$.

3.3 Computing Utility

We now define *Utility* : $\mathcal{T}_A \mapsto [0, 1]$ of a pivot table T by combining *Insightfulness* (Eq. 5) and *Interpretability* (Eq. 9). To balance their contributions, we introduce a tunable parameter α , set to 0.5 by default to give equal weight to both components. However, α can be adjusted to reflect application-specific preferences.

$\text{Utility}(T) = \alpha \cdot \text{Insightfulness}(T) + (1 - \alpha) \cdot \text{Interpretability}(T)$

EXAMPLE 3.11. For Figure 5(d), the *Insightfulness* and *Interpretability* scores are 0.39 and 0.94, respectively. Thus, $\text{Utility}(T) = 0.5 \times 0.39 + 0.5 \times 0.94 = 0.67$.

4 Diversity in a Set of Pivot Tables

While utility quantifies the goodness of a single pivot table in isolation, *diversity* captures how well a *set* of pivot tables, *collectively*, provide complementary and unique perspectives on the data (D3). High diversity in a set of pivot tables is achieved when the pivot tables are distant from each other with respect to data coverage and the insights they provide.

Diversity. Following Max-Min diversification [2], we define diversity of a set of pivot tables $T = \{T_1, T_2, \dots\}$ by the smallest pairwise distance between T 's elements. More formally, given a symmetric distance function $\text{dist} : \mathcal{T}_A \times \mathcal{T}_A \mapsto [0, 1]$, we define *diversity* : $2^{\mathcal{T}_A} \mapsto [0, 1]$ of a set of pivot tables $T \subseteq \mathcal{T}_A$ as follows:

$$\text{Diversity}(T) = \min_{T_i, T_j \in T \text{ s.t. } i < j} \text{dist}(T_i, T_j)$$

Distance between pivot tables. A simple heuristic to model *dist* is the degree of disjointness between the attributes that define the pivot-table queries—if two pivot tables operate on the same set of attributes, the distance is 0; for completely disjoint set of attributes, the distance is 1. However, this heuristic fails to account for the *structural semantics* of the pivot-table queries and the *content semantics* of the data in the pivot tables. To this end, we employ a semantics-preserving embedding function $E : \mathcal{T}_A \mapsto [-1, 1]^p$, which maps a pivot table $T \in \mathcal{T}_A$ to a p -dimensional vector. Then, we compute distances between pivot tables in this embedding space:

$$\text{dist}(T_1, T_2) = \frac{1 - \text{cosine_similarity}(E(T_1), E(T_2))}{2}$$

Here, $\text{cosine_similarity} : \mathbb{R}^p \times \mathbb{R}^p \rightarrow [-1, 1]$ is a widely used measure for comparing embeddings [5, 54]. We divide by 2 to achieve normalization s.t. $\text{dist}(T_1, T_2) \in [0, 1]$.

Pivot-table embedding. Pivot-table embedding should capture both the syntactic and semantic characteristics of the pivot table. To this end, we combine the query embedding E_Q and the content embedding E_C through concatenation, i.e., $E(T) = [E_Q(T); E_C(T)]$. Concatenating embeddings is a widely used technique in machine learning, natural language processing, and multi-modal learning [69].

Query embedding. Queries define the structural intent of a pivot table. While they reference data attributes and are aware of the schema, they are agnostic to the pivot table's content. Thus, the same pivot-table query over different database instances with an identical schema should yield the same query embedding. Effective query embeddings must also capture the semantics of the attribute names and reflect the query semantics. E.g., `GROUP BY Income` and `GROUP BY Salary` are semantically similar and should therefore be close in the query-embedding space. To this end, we use T5 [53], a natural-language encoder fine-tuned over a Text-to-SQL dataset [68], to obtain the query embedding $E_Q : \mathcal{T}_A \mapsto [-1, 1]^{1024}$.

Content embedding. The content embedding must capture both the statistical and distributional properties of the pivot-table data and structural relationships among its attributes and tuples. In this work, we leverage TAPEX [40], a pre-trained encoder trained on sentence-table pairs, which is designed to understand both the structure and content of tabular data. This results in a content embedding $E_C : \mathcal{T}_A \mapsto [-1, 1]^{1024}$.

5 The SAGE Algorithm

We now present a solution to Problem 2.1, which recommends a k -budgeted set of pivot tables with maximum utility under a diversity constraint. This is an instance of the NP-hard Maximum Weight Independent Set (MWIS) problem [56], whose exact solution requires an exponential time, with complexity $O(|\mathcal{T}_A|^k)$ in the size of the search space (\mathcal{T}_A in our case). Furthermore, \mathcal{T}_A grows combinatorially with $|A|$, leading to an exponential growth. Approximation techniques exist for special cases of MWIS and related problems—such as sequential [56] and distributed [36] greedy algorithms, genetic algorithms [49], and integer linear programming [12]. Alternative formulations are possible such as incorporating diversity into the objective function or solving the dual version that maximizes diversity [2] under a utility constraint.

However, in SAGE, we adopt a simple greedy approach (Section 5.2) for two key reasons: (1) Interactive response time is desirable for our problem settings and a greedy approach achieves linear time complexity of $O(|\mathcal{T}_A|)$. (2) As we demonstrate in Section 6, our greedy approach works remarkably well in practice over real-world datasets, almost always matching the exact solution.

5.1 Optimizations

The linear-time complexity of the greedy approach is still prohibitive for practical use for two reasons: (1) Computing the *Utility* of candidate pivot tables requires their materialization, which is computationally expensive—especially since \mathcal{T}_A grows exponentially with A . (2) Some components of *Utility* relies on the LLM, whose inference latency is a bottleneck due to its very large number of parameters [61, 63]. To address these efficiency challenges (D5), we introduce two offline optimizations. The first is *candidate pruning*, solely based on the query structure of the pivot tables, to eliminate potentially low-utility ones before materialization. Pruning significantly reduces the search space and avoids a large number of materializations (§ 5.1.1). The second is a lightweight *proxy model*, tailored to the dataset, to approximate LLM inferences efficiently (§ 5.1.2).

5.1.1 Pruning. To prune pivot tables likely to yield low utility without materialization, we leverage three components of *Utility* (§ 3.3): C-I: *attribute significance* (§ 3.1.1) used in *Insightfulness* (Eq 5), C-II: *semantic validity* (§ 3.2.2), and C-III: *conciseness* (§ 3.2.3) used in *Interpretability* (Eq 9). These choices are motivated by computational efficiency, as these components require knowledge of only the pivot-table query, which defines the structure of the pivot table—such as which attributes are used, number of cells, etc.—and do not require a full materialization over the dataset.

While *Insightfulness* selects only the maximum among multiple components, *Interpretability* averages out three components—two of which are C-II and C-III above. This enables effective pruning of pivot tables that already show low scores for C-II and/or C-III. Furthermore, since C-I interacts in a multiplicative way with other components of *Insightfulness* (Eq 5), a low value will inevitably result in low *Insightfulness*, making it an ideal choice. Our pruning algorithm works as follows: we evaluate scores for the above three components and over-approximate *Utility* (§ 3.3) by assigning maximum possible values for all other components whose values are unknown. Through this conservative estimate, we discard candidates with a score below a threshold (a system parameter set to 0.5).

Algorithm 1: SAGE algorithm

Input : Database D ,
 Unmaterialized candidate pivot table set \mathcal{T}_A ,
 Diversity threshold θ ,
 The desired number of pivot tables k

Output : A high-utility set $T \subseteq \mathcal{T}_A$ s.t. $|T| \leq k$ and $Diversity(T) \geq \theta$

```

/* Offline phase: executed once for each dataset. Re-executed if the
  schema changes or major change happens in the data distribution. */
/* Prune candidate set based on pivot-table structure (Section 5.1.1) */
1  $\mathcal{T}_A^{Pr} \leftarrow Prune(\mathcal{T}_A)$ 
/* Building LLM-Proxy (Section 5.1.2) */
2  $Q \leftarrow \text{list of questions about } D$  /* Generate prompts for the LLM */
3  $R \leftarrow \text{LLM-Response}(Q)$  /* Get responses from the LLM */
4  $LLM\text{-Proxy} \leftarrow Train(Q, R)$  /* Train a proxy prediction model */

/* Online phase: executed whenever the data changes. */
/* Materialize and compute utility based on pivot-table contents */
5 foreach  $T \in \mathcal{T}_A^{Pr}$  do
6   Materialize  $T$  over  $D$ 
7   Compute  $Utility(T)$  /* Section 3.3 */
8  $\mathcal{T}_A^{Pr} = sorted(\mathcal{T}_A^{Pr})$  /* Sort by the descending order of Utility */
/* Greedy selection (Section 5.2) */
9  $T \leftarrow \emptyset$  /* Initialize an empty set */
10 while  $|T| \leq k$  do
11   foreach  $T \in \mathcal{T}_A^{Pr}$  do
12     /* No pivot table in  $T$  is within  $\theta$  distance away from  $T$  */
13     if  $\nexists T' \in T$  s.t.  $dist(T, T') < \theta$  then
14        $T \leftarrow T \cup \{T\}$ 
14 return  $T$ 

```

We note that data sampling, approximate query processing [3, 15, 38], and incorporating domain-specific insights are other avenues for pruning. While we limit SAGE to query-structure-based pruning, other pruning techniques can be incorporated in SAGE.

5.1.2 LLM-proxy. Recall that the computation of pivot-table utility requires LLM inferences (Section 3), and each query to LLM is time-consuming in practice. Even if we cache the LLM responses, any change in the underlying data or user-specified parameters (D4) would require LLM re-consultation. To expedite this process, we train a cheap but significantly faster decision-tree classifier to mimic LLM behavior, serving as an “LLM-proxy” during the online phase of SAGE (Algorithm 1). We train the classifier on LLM prompt-response pairs, where we generate potential LLM-queries based on the dataset. The proxy model is a simple prediction model to answer the fixed-template questions discussed in Section 3. Note that we do not need to re-train this proxy model as long as the data distribution and overall trends in the dataset remain unchanged. We empirically found these LLM-proxies to achieve about 90% accuracy for trend-related predictions (§ 3.1.3) and 60% accuracy for surprise-related predictions (§ 3.1.4). While the accuracy can further be improved with additional training or using more complex classifiers, these accuracies are sufficient for practical applications. The key benefit here is that replacing expensive LLM calls with fast proxy predictions can substantially expedite the online recommendation phase.

5.2 Greedy Algorithm

Algorithm 1 shows the SAGE workflow. Line 1 denotes the offline pruning step and lines 2–4 show the steps for LLM-proxy training.

The online phase is shown in lines 5–14. The pivot tables that survive the pruning step are materialized in line 6, we then compute their *Utility* scores in line 7, and sort the pruned candidate set of pivot tables by descending order of utility in line 8.

The greedy selection phase is shown in lines 9–14, where we select pivot tables greedily while ensuring that they satisfy the diversity constraint w.r.t the already selected ones in T . If a candidate pivot table is at least θ away from all the previously selected tables, we include it to T (lines 12–13). The algorithm terminates when we have selected k pivot tables or when no more candidates remain that satisfy the diversity constraint.

Tuning the diversity threshold θ . When the user is unable to specify or tune a desired diversity threshold θ , either due to unfamiliarity of the dataset or lack of expertise, the desired number of pivot tables k can be used as a guideline to derive the value of θ . In that scenario, a trivial extension to our approach would be to cluster the candidate pivot tables into k groups, and then apply a greedy strategy to maximize a linear combination of *Utility* (Section 3) and *Diversity* (Section 4) as was done in prior work [31].

6 Experimental Results

We now present experimental results to demonstrate the efficacy of SAGE in practical settings to address the following questions:

- (Q1) How do SAGE’s runtime and recommendation quality compare quantitatively with those of existing methods? (§ 6.2)
- (Q2) What is the effect of the optimization techniques—pruning and LLM-proxy—on SAGE’s runtime performance? (§ 6.3)
- (Q3) How well does SAGE scale with data growth? (§ 6.4)
- (Q4) How do key parameters (budget k and diversity threshold θ) influence the quality of SAGE recommendations? (§ 6.5)
- (Q5) How do SAGE recommendations qualitatively compare against commercial software and LLMs over real datasets? (§ 6.6)

6.1 Experimental Setup

All experiments were run on machines with 256 GB RAM running Ubuntu 22.04 LTS with CPU 12 cores and GPU NVIDIA H100 96GB with CUDA 12.8. We implemented our solutions with Python 3.10.3, leveraging the Pandas library for data manipulation and pivot table generation. For embeddings, we utilized TAPEX-large [43] and T5 trained by Spider [26]. We employed Llama-3-7B [41] as the LLM for semantic consultation. Our source code is publicly available [6].

6.1.1 Datasets. We used three real-world datasets in our experiments. These datasets vary in domain and size, allowing us to assess SAGE’s generalizability across different contexts.

Marketing. The Marketing dataset [55] comprises demographic and behavioral information about customers, including marital status and purchase history. It consists of 2,240 tuples and 28 attributes, encompassing a diverse set of features: 9 numerical attributes (e.g., Year_Birth, Income) and 19 categorical attributes (e.g., Education, Marital_Status, Complaint).

Video. The Video dataset [29] comprises video game sales from various countries, platforms, or release years. It contains 16,600 tuples across 11 attributes. Out of these, 7 attributes are categorical (e.g., Platform, Genre, Publisher) and 4 are numerical (e.g., Rank, Year, NA_Sales, Global_Sales).

		Marketing							Video							House						
		#PT	T(s)	Ins	Int	Util	m-dist	Div	#PT	T(s)	Ins	Int	Util	m-dist	Div	#PT	T(s)	Ins	Int	Util	m-dist	Div
$k=3, \theta=0.3$ or 0.2	Top-k	3	147	2.86	2.23	2.55	0.39		3	64	1.99	1.27	1.63	0.06		3	2	2.26	2.13	2.20	0.09	
	DAISY	3	23	0.54	1.48	1.01	0.16		3	21	1.63	1.00	<u>1.32</u>	0.04		3	455	0.23	1.17	0.70	0.44	
	LLMs	3	15	0.16	1.75	0.96	0.16		3	25	0.01	0.58	0.29	<u>0.32</u>		3	30	0.07	0.55	0.31	0.06	
	PowerBI	3	9	0.91	0.94	0.93	0.41		3	6	0.91	0.94	0.93	0.36		3	15	0.50	1.76	1.13	<u>0.36</u>	
	GSheets	2	1	0.00	1.67	0.83	0.56		3	1	0.04	0.59	0.31	<u>0.32</u>		–	–	–	–	–	–	–
	Excel	3	1	0.29	2.42	1.35	0.14		3	1	0.14	2.16	1.15	0.21		3	1	0.00	1.87	0.93	0.31	
	SAGE	3	161	2.86	2.23	2.55	<u>0.45</u>		3	69	1.99	1.27	1.63	0.30		3	2	2.21	2.09	<u>2.15</u>	0.21	
$k=5, \theta=0.3$ or 0.2	Top-k	5	148	4.76	3.72	4.24	0.11		5	64	3.31	2.03	2.67	<u>0.05</u>		5	2	3.77	3.56	3.66	0.05	
	DAISY	5	23	0.54	2.51	1.52	0.16		5	21	2.57	1.67	<u>2.21</u>	0.04		5	455	0.23	1.84	1.04	<u>0.29</u>	
	LLMs	5	22	0.93	2.94	1.94	<u>0.31</u>		5	20	0.18	0.86	0.52	0.03		5	44	0.07	0.75	0.41	0.13	
	PowerBI	5	9	0.20	4.23	2.21	0.25		–	–	–	–	–	–		5	15	0.50	2.99	1.74	0.36	
	Excel	5	1	0.29	3.98	2.13	0.08		5	1	0.31	3.28	1.80	0.04		5	1	0.00	3.05	1.52	0.08	
	SAGE	5	161	4.76	3.70	<u>4.23</u>	0.33		5	69	3.31	2.03	2.67	0.30		5	2	3.42	3.52	<u>3.47</u>	0.20	
	Top-k	10	147	9.53	7.44	8.49	0.02		10	64	4.06	5.50	4.78	<u>0.05</u>		10	2	7.57	7.02	7.30	0.05	
$k=10, \theta=0.1$	DAISY	10	23	0.54	4.22	2.38	0.12		10	21	4.09	3.33	<u>3.71</u>	0.02		10	455	0.23	3.28	1.75	<u>0.11</u>	
	LLMs	10	32	1.70	5.67	3.69	0.20		10	28	0.74	2.24	1.49	0.04		10	43	0.26	1.72	0.99	0.02	
	PowerBI	9	9	0.20	4.88	2.54	0.09		–	–	–	–	–	–		10	15	0.50	5.99	3.24	0.14	
	Excel	7	1	0.29	5.30	2.79	0.08		9	1	0.31	4.71	2.51	0.01		7	1	0.00	3.97	1.98	0.08	
	SAGE	10	161	5.58	9.53	<u>7.44</u>	<u>0.11</u>		10	69	4.06	5.50	4.78	0.11		10	2	7.55	6.91	<u>7.23</u>	0.10	

Table 1: Comparison with baselines across three datasets and three values of k . For $k = 3$ and $k = 5$, we used $\theta = 0.3$ on Marketing and Video, and $\theta = 0.2$ on House. For $k = 10$, $\theta = 0.1$ was used across all datasets. Columns show the number of pivot tables recommended (#PT), elapsed time (T, in seconds), total *Insightfulness* (Ins), *Interpretability* (Int), *Utility* (Util), and *Diversity* in terms of minimum pairwise distance (m-dist) and a distance matrix visualization (Div). In the distance matrix, the diagonal is black, denoting 0 self distance from a pivot table to itself. Lighter colors denote less similar pairs, offering diversity. For Excel, Power BI, and Google Sheets, k cannot not be controlled. Therefore, we consider the first- k items when more than k are returned. The best values in Util and m-dist are marked as bold and the second best underlined.

House. The House dataset [4] contains information about residential property sales, comprising 1,460 tuples and 81 attributes. Among the attributes, 23 are numerical (e.g., SalePrice, LotArea, OverallQual, GrLivArea) and 58 are categorical (e.g., Street, HouseStyle, RoofStyle, GarageType).

6.1.2 Baselines. We consider the following baselines:

Brute-Force. We employ an exhaustive Brute-Force search, which considers all possible k -sized pivot table sets and follows a naive exhaustive approach to solve Problem 2.1. Note that in the absence of an absolute ground-truth, the results obtained from this technique can be treated as the optimal solution.

Top-k. This baseline ranks candidate pivot tables in descending order of their utility scores and selects the top- k as the recommended set, without accounting for the diversity constraint. To ensure a fair comparison, we apply our optimization techniques (pruning and LLM-proxy) prior to the selection phase. This allows it to generate recommendations within a runtime comparable to that of SAGE.

LLM. We used Llama-3-8B-Instruct (Meta) [41] as a transformer-based large language model (LLM) as another baseline. While LLMs are increasingly used for tabular data tasks, few tools address our specific problem. We prompted the model with a small data sample and asked for interesting and diversified pivot tables in natural language, according to our problem setup.

DAISY. DAISY [65] is a query recommendation system trained on “interesting” queries. Due to the original model and data being publicly unavailable, we put our best effort to replicate a DAISY-like model. We generated training data from Auto-Suggest pivot tables [66], assuming the pivot tables suggested by Auto-Suggest are insightful. We created less-interesting (negative) examples by replacing attributes in these insightful tables with random ones, and then trained a classifier to distinguish them.

Microsoft Excel. Microsoft Excel [44] is a widely used commercial spreadsheet software that offers built-in pivot table recommendations based on the underlying data. We used the Windows version 2501, evaluated during February 2025.

PowerBI. Microsoft PowerBI [21, 45] is a business intelligence software that offers “quick insights” [21] in various forms including visualizations. For comparison, we focus only on the recommendations where the underlying query excludes the WHERE clause, which aligns with our settings.

Google Sheets. Google Sheets [28] is an online spreadsheet software known for easy collaboration. Google Sheets provide automatic recommendation of pivot tables for a dataset. We used the browser version during the month of February, 2025.

For Excel, PowerBI, and Google Sheets, k cannot be controlled. Thus, we consider the first- k items when more than k were returned.

6.2 Contrasting against Baselines

Table 1 contrasts SAGE against the baselines across three datasets. Our primary metrics for comparison are *Utility* (Util) and *Diversity* (m-dist), but we report *Insightfulness* and *Interpretability* for additional context. The results show that SAGE effectively balances utility and diversity, outperforming approaches like Top-k, which inherently lack diversity. For Marketing, SAGE achieves a strong overall performance, ranking either best or second-best in both metrics. Notably, when SAGE marginally loses in terms of diversity, that is usually complemented by significantly higher utility. For example, in Marketing, $k = 10$, LLM’s diversity (0.20) is more than SAGE’s diversity (0.11). However, LLM’s utility (3.69) is about half of SAGE’s utility (7.44). A similar situation is seen for $k = 3$ where GSheets yields only 30% of SAGE’s utility. Recall that SAGE’s goal is to just satisfy the diversity constraint, not maximize it.

For Video, SAGE outperforms all baselines across all cases in terms of utility and two cases in terms of diversity. SAGE and Top-k’s similar performance can be attributed to the dataset attributes, such as North-America Sales and Europe-Sales, which possess similar value ranges, naturally leading to comparable utility different selections. For House, Top-k marginally outperforms SAGE in terms of utility, but at the cost very low diversity. SAGE could not offer more diversity here, because the pruning phase retained only 11 of 81 attributes, where other baselines used all attributes.

While Top-k achieves high utility by design, it fails to diversify. LLM performs well on Marketing but struggles with Video and House, because Marketing has many categorical values, which LLMs are adept at interpreting and utilizing for diverse recommendations. While PowerBI demonstrates good diversity across the board, it typically yields poor utility. Google Sheets demonstrates good diversity on Marketing and Video, however, its recommendations are limited to a small number of tables, leading to low utility. Google Sheets also fails to recommend for House, likely due to high dimensionality (81 attributes), indicating its limitation in handling complex, high-dimensional datasets. Excel consistently shows reasonable utility, however, it underperforms in terms of diversity, probably due to its tendency to allow significant column overlaps.

6.3 Effect of Optimizations

Table 2 demonstrates how the optimizations—pruning (PR) and LLM-proxy (PX)—significantly improve SAGE’s runtime performance while costing minimal utility loss. For this experiment, we used a vertical slice of the Marketing dataset over 5 attributes to allow Brute Force to finish within a reasonable time. We performed an ablation study where we turned off all the optimizations (No PR/PX = plain greedy, no pruning (No PR), no LLM-proxy (No PX), and SAGE with both optimizations. Notably, we find that greedy achieves exact results in par with Brute Force, validating our choice.

LLM-proxy causes slight decrease (3%) in utility due to less accurate proxy classifier, however, boosts performance significantly. Pruning does not hurt utility, due to its conservative filtering of unpromising candidates. Figure 9 (left) shows runtime comparison over three datasets (limited to 5 attributes), where both optimizations offer significant performance boost across the board. Pruning is significantly beneficial for high-dimensional datasets, due to the exponential growth of the candidate space w.r.t attributes.

	Approach	BruteForce	No PR/PX	No PR	No PX	SAGE
	#Pivot Tables (PTs)	130	130	130	20	20
$K = 2$	#PT combinations	8,385	N/A	N/A	N/A	N/A
	Runtime (s)	5,817	4,971	76	492	23
	Utility (%)	100	100	97	100	97
	– Insightfulness (%)	100	100	94	100	94
	– Interpretability (%)	100	100	100	100	100
$K = 5$	#PT combinations	286M	N/A	N/A	N/A	N/A
	Runtime (s)	19,604	4,971	76	492	24
	Utility (%)	100	100	97	100	97
	– Insightfulness (%)	100	100	92	100	92
	– Interpretability (%)	100	100	100	100	100
PR = Pruning optimization		No PR/PX = Plain greedy, no optimization				
PX = LLM-Proxy optimization		SAGE = Greedy + PR + PX				

Table 2: Comparison among variants for $\theta = 0.1$ on the Marketing dataset over 5 attributes. We report the runtimes for the online phase for SAGE and other variants when optimizations are applied.




Figure 8: SAGE runtime (s) w.r.t (Left) #tuples, (Right) #attributes. We used $k = 5$ and $\theta = 0.1$ for these experiments.

6.4 Scalability

Figure 8 shows the runtime of SAGE w.r.t the number of tuples (left) and the number of attributes (right), averaged over three executions. We find that SAGE’s runtime increases linearly with larger data volumes and quadratically with higher dimensionality, because increasing dimensions significantly expands the combinatorial search space. The Marketing dataset exhibits the highest runtime due to its high dimensionality, while House remains consistently fast. The pruning mechanism filters out most attributes in Video and House, retaining only 7 and 12 attributes, respectively, where Marketing retains 17 attributes. Marketing incurs a heavier computational burden due to its varied numerical ranges and high dimensionality. For House, the execution time is minimal because SAGE prunes most attributes, reducing the runtime to just a few milliseconds.

In summary, SAGE scales linearly as data grows vertically (#tuples) and quadratically as data grows horizontally (#attributes), while its runtime stays within practical range. Stricter pruning or sampling can be used to further boost runtime performance.

6.5 Parameter Sensitivity

Figure 9 (right) depicts the impact of varying the diversity threshold (θ) on the utility score over the Video dataset. We observe that as θ increases, the utility score drops for all k . This is expected because a higher θ constrains the selection of candidates more stringently, leading to fewer eligible items. To satisfy this stricter criterion, the algorithm may be compelled to select items with lower utility to maintain diversity, thus decreasing the total utility. The utility drop

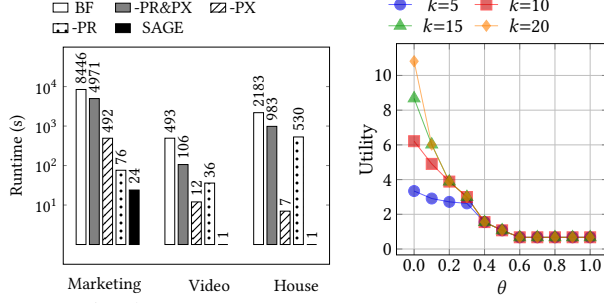


Figure 9: (Left) Effect of optimization techniques. We used $k=4$ and $\theta = 0.1$. (Right) Effect of θ on Utility on the Video dataset.

becomes particularly significant for high values of k , as satisfying a high diversity threshold for a larger result set is more challenging.

6.6 Case Studies

We now present findings from case studies in which we manually examined the pivot tables for qualitative insights regarding recommendations made by SAGE vs. commercial software and LLMs.

Diverse and meaningful aggregates. We found SAGE to consistently recommend diverse and semantically meaningful aggregate functions. For the Video dataset, SAGE recommends pivot tables that include aggregates such as COUNT, MAX, MIN, SUM, and MEAN. In contrast, Excel, Google Sheets, and PowerBI typically involve only SUM or COUNT. While LLMs too suggest a variety of aggregate functions, they sometimes generate hallucinated attributes. For example, the attribute `LotLocation` is not present in the House dataset, but the LLM recommends an aggregate involving it. Furthermore, despite requesting to use the specified aggregate functions, LLM selects out-of-scope aggregates such as median or standard deviation. Generally, these tools often fail to provide semantics-aware recommendations, for example, Excel once suggested `SUM(Birth_Year)`, a completely meaningless aggregate.

Diverse and meaningful attributes. We found SAGE to consistently select diverse and semantically meaningful attributes for aggregation and grouping. In the Video Dataset, SAGE utilizes Publisher, Sales, and Year for various aggregations, providing comprehensive coverage of the attribute space. This contrasts sharply with Excel and Google Sheets, which make repetitive GROUPBY choices, while LLM consistently focuses only on Sales for aggregation. In the House Dataset, SAGE identifies attributes such as `KitchenAbvGr`, `HalfBath`, and `FullBath` for aggregations that other baselines completely ignore. While only SAGE and LLM appropriately use `SalesPrice` for aggregation, LLM limits itself to aggregate only over this attribute. In contrast, SAGE diversifies selection of attributes, avoiding meaningless choices made by other tools, such as PowerBI’s poor recommendation to GROUP BY ID, which results in an extremely large and incomprehensible summary.

Avoiding unsurprising summaries. Since SAGE can assess the degree of surprisingness in observed trends, it avoids recommending obvious or trivial queries that offer little utility. Unlike other tools that consistently recommend common aggregations such as `MEAN(INCOME)`, regardless of context, SAGE avoids recommendations that provide obvious insights (e.g., people with low income

spend less). For the Video dataset, existing baselines typically focused on aggregations such as `SUM(SALES)`. In contrast, SAGE included diverse types of aggregations such as `MEAN(NA_Sales)` and `MAX(JP_Sales)`. Additionally, it discovered more surprising patterns, such as `COUNT(YEAR) GROUP BY GENRE`, which were not considered by others.

7 Related Work

Data Summarization. Prior work defines informative summaries as tables that exhibit significant discrepancies between groups [17, 30, 62]. Tools in this domain aim at identifying interesting patterns using deviation scores [62], correlation, and outlier detection [17, 21, 30]. Smart Drill-Down [35] enable users to discover and summarize interesting groups of tuples described by rules. Other works [58, 59] explore user-adaptive exploration by guiding users to the surprising data regions, or characterizing informative content based on user familiarity. However, these approaches primarily prioritize identifying best summaries based on a specific metric, often leading to top-k recommendations, and do not consider diversity.

History-based Recommendation. Auto-Suggest [66] and DAISY [65] leverage large-scale user logs, SQL queries, or crowdsourcing to predict the most appropriate pivot table for a given database. While these systems can generate recommendations based on historical usage patterns, they disregard the aspect of diversity. Moreover, relying on historical data cannot guarantee table informativeness since it does not validate the generated pivot tables.

Diversification Algorithms. In conventional query result recommendation systems, diversity definitions are typically categorized into three types: (1) content-based diversity, (2) novelty, and (3) coverage [23, 24]. Max-Sum diversification [13, 23, 52] maximizes the linear combination of diversity and utility scores, while Max-Min diversification maximizes the minimum diversity between selected items. DisC [24] computes diversity by evaluating data equivalence across dimensions. However, existing diversification methods are designed for tuples or documents rather than entire tables as in our case. The applicability of these algorithms to table-based recommendation has not been explored in prior research.

8 Conclusions and Future Work

We presented SAGE to recommend diverse set of pivot tables while balancing insightfulness and interpretability. We introduced a utility model for a single pivot table, a diversity metric for pivot-table sets, and a simple greedy algorithm built on two optimization techniques for efficient recommendation. We empirically showed that SAGE outperforms baselines in diversity while maintaining high utility. Our case studies illustrated SAGE’s ability to avoid generic patterns and provide data-semantics-aware suggestions. To the best of our knowledge, this is the first work to combine diversity and data-semantics for data summarization.

SAGE currently does not incorporate contextual information, such as the user’s workflow (e.g., their end goals) or broader ecosystem (e.g., tools in their software stack). Integrating such context could significantly enhance the quality of recommendations. Another promising direction is to support alternative forms of data summaries—such as textual descriptions that highlight key trends—in addition to structured formats like pivot tables.

References

- [1] 2023. Microsoft Community Forum: Issues with null values in pivot tables. <https://answers.microsoft.com/en-us/msoffice/forum/all/how-to-get-rid-of-blank-appearing-in-pivot-table/449f785a-f993-4c40-922b-1b52f02571ce>. Accessed: 2025-06-02.
- [2] Raghavendra Addanki, Andrew McGregor, Alexandra Meliou, and Zafeiria Mounmolidou. 2022. Improved Approximation and Scalability for Fair Max-Min Diversification. In *25th International Conference on Database Theory, ICDT 2022, March 29 to April 1, 2022, Edinburgh, UK (Virtual Conference) (LIPIcs, Vol. 220)*, Dan Olteanu and Nils Vortmeier (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 7:1–7:21. doi:10.4230/LIPIcs.ICDT.2022.7
- [3] Sameet Agarwal, Rakesh Agrawal, Prasad Deshpande, Ashish Gupta, Jeffrey F. Naughton, Raghu Ramakrishnan, and Sunita Sarawagi. 1996. On the Computation of Multidimensional Aggregates. In *VLDB'96, Proceedings of 22th International Conference on Very Large Data Bases, September 3-6, 1996, Mumbai (Bombay), India*, T. M. Vijayaraman, Alejandro P. Buchmann, C. Mohan, and Nandlal L. Sarda (Eds.). Morgan Kaufmann, 506–521. <http://www.vldb.org/conf/1996/P506.PDF>
- [4] animeshparikshya. [n.d.]. House Sale Data (81 Column) Dataset. Kaggle Dataset. <https://www.kaggle.com/datasets/animeshparikshya/house-sale-data-81-colum> Accessed 16 July 2025.
- [5] Sanjeev Arora, Yingyu Liang, and Tengyu Ma. 2017. A Simple but Tough-to-Beat Baseline for Sentence Embeddings. In *5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24-26, 2017, Conference Track Proceedings*. OpenReview.net. <https://openreview.net/forum?id=SyK00v5xx>
- [6] Anonymous Authors. 2025. SAGE: Data-Semantics-Aware Recommendation of Diverse Pivot Tables. <https://anonymous.4open.science/r/SAGE-BE88/>.
- [7] Pierre Barroillet, Sophie Bernardin, Sophie Portrat, Evie Vergauwe, and Valérie Camos. 2007. Time and cognitive load in working memory. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 33, 3 (2007), 570.
- [8] Leilani Battle and Jeffrey Heer. 2019. Characterizing Exploratory Visual Analysis: A Literature Review and Evaluation of Analytic Provenance in Tableau. *Comput. Graph. Forum* 38, 3 (2019), 145–159. doi:10.1111/CGF.13678
- [9] P. J. Bickel, E. A. Hammel, and J. W. O'Connell. 1975. Sex Bias in Graduate Admissions: Data from Berkeley. *Science* 187, 4175 (1975), 398–404. doi:10.1126/science.187.4175.398 arXiv:https://www.science.org/doi/pdf/10.1126/science.187.4175.398
- [10] Colin R. Blyth. 1972. On Simpson's Paradox and the Sure-Thing Principle. *J. Amer. Statist. Assoc.* 67, 338 (1972), 364–366. <http://www.jstor.org/stable/2284382>
- [11] Allan Borodin, Hyun Chul Lee, and Yuli Ye. 2012. Max-Sum diversification, monotone submodular functions and dynamic updates. In *Proceedings of the 31st ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems, PODS 2012, Scottsdale, AZ, USA, May 20-24, 2012*, Michael Benedikt, Markus Krötzsch, and Maurizio Lenzerini (Eds.). ACM, 155–166. doi:10.1145/2213556.2213580
- [12] Matteo Brucato, Juan Felipe Beltran, Azza Abouzied, and Alexandra Meliou. 2016. Scalable Package Queries in Relational Database Systems. *Proc. VLDB Endow.* 9, 7 (2016), 576–587. doi:10.14778/2904483.2904489
- [13] Jaime G. Carbonell and Jade Goldstein. 1998. The Use of MMR, Diversity-Based Reranking for Reordering Documents and Producing Summaries. In *SIGIR '98: Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, August 24-28 1998, Melbourne, Australia*, W. Bruce Croft, Alistair Moffat, C. J. van Rijsbergen, Ross Wilkinson, and Justin Zobel (Eds.). ACM, 335–336. doi:10.1145/290941.291025
- [14] Minmin Chen, Yuyan Wang, Can Xu, Ya Le, Mohit Sharma, Lee Richardson, Su-Lin Wu, and Ed H. Chi. 2021. Values of User Exploration in Recommender Systems. In *RecSys '21: Fifteenth ACM Conference on Recommender Systems, Amsterdam, The Netherlands, 27 September 2021 - 1 October 2021*, Humberto Jesús Corona Pampin, Martha A. Larson, Martijn C. Willemsen, Joseph A. Konstan, Julian J. McAuley, Jean Garcia-Gathright, Bouke Huurnink, and Even Oldridge (Eds.). ACM, 85–95. doi:10.1145/3460231.3474236
- [15] Graham Cormode and S. Muthukrishnan. 2005. An improved data stream summary: the count-min sketch and its applications. *J. Algorithms* 55, 1 (2005), 58–75. doi:10.1016/J.JALGOR.2003.12.001
- [16] Anna DeCastellarnau. 2018. A classification of response scale characteristics that affect data quality: a literature review. *Quality & Quantity* 52, 4 (2018), 1523–1559.
- [17] Çagatay Demiralp, Peter J. Haas, Srinivasan Parthasarathy, and Tejaswini Pedapati. 2017. Foresight: Recommending Visual Insights. *Proc. VLDB Endow.* 10, 12 (2017), 1937–1940. doi:10.14778/3137765.3137813
- [18] Dazhen Deng, Aoyu Wu, Huamin Qu, and Yingcai Wu. 2023. DashBot: Insight-Driven Dashboard Generation Based on Deep Reinforcement Learning. *IEEE Trans. Vis. Comput. Graph.* 29, 1 (2023), 690–700. doi:10.1109/TVCG.2022.3209468
- [19] Mukund Deshpande and George Karypis. 2004. Item-based top-n recommendation algorithms. *ACM Transactions on Information Systems (TOIS)* 22, 1 (2004), 143–177.
- [20] Diginomica. 2025. *Shadow IT Never Dies: Why Spreadsheets Are Still Running Your Business*. <https://diginomica.com/shadow-it-never-dies-why-spreadsheets-are-still-running-your-business> Accessed: 2025-01-21.
- [21] Rui Ding, Shi Han, Yong Xu, Haidong Zhang, and Dongmei Zhang. 2019. Quick-Insights: Quick and Automatic Discovery of Insights from Multi-Dimensional Data. In *Proceedings of the 2019 International Conference on Management of Data, SIGMOD Conference 2019, Amsterdam, The Netherlands, June 30 - July 5, 2019*, Peter A. Boncz, Stefan Manegold, Anastasia Ailamaki, Amol Deshpande, and Tim Kraska (Eds.). ACM, 317–332. doi:10.1145/3299869.3314037
- [22] Marina Drosou, H. V. Jagadish, Evaggelia Pitoura, and Julia Stoyanovich. 2017. Diversity in Big Data: A Review. *Big Data* 5, 2 (2017), 73–84. doi:10.1089/BIG.2016.0054
- [23] Marina Drosou and Evaggelia Pitoura. 2010. Search result diversification. *SIGMOD Rec.* 39, 1 (2010), 41–47. doi:10.1145/1860702.1860709
- [24] Marina Drosou and Evaggelia Pitoura. 2012. DisC diversity: result diversification based on dissimilarity and coverage. *Proc. VLDB Endow.* 6, 1 (2012), 13–24. doi:10.14778/2428536.2428538
- [25] Anna Fariha, Ashish Tiwari, Arjun Radhakrishna, Sumit Gulwani, and Alexandra Meliou. 2021. Conformance Constraint Discovery: Measuring Trust in Data-Driven Systems. In *SIGMOD '21: International Conference on Management of Data, Virtual Event, China, June 20-25, 2021*, Guoliang Li, Zhanhuai Li, Stratos Idreos, and Divesh Srivastava (Eds.). ACM, 499–512. doi:10.1145/3448016.3452795
- [26] GaussAlgo. [n.d.]. T5-LM-Large-text2sql-spider Model Card. Hugging Face Model Card. <https://huggingface.co/gaussalgo/T5-LM-Large-text2sql-spider> Accessed 16 July 2025.
- [27] Kareem El Gebaly, Parag Agrawal, Lukasz Golab, Flip Korn, and Divesh Srivastava. 2014. Interpretable and Informative Explanations of Outcomes. *PVLDB* 8, 1 (2014), 61–72. doi:10.14778/2735461.2735467
- [28] Google. 2025. Google Sheets - Online Spreadsheet Editor. <https://sheets.google.com/> Accessed: 2025-01-21.
- [29] gregorut. [n.d.]. Video Game Sales Dataset. Kaggle Dataset. <https://www.kaggle.com/datasets/gregorut/videogamesales> Accessed 16 July 2025.
- [30] Camille Harris, Ryan A. Rossi, Sana Malik, Jane Hoffswell, Fan Du, Tak Yeon Lee, Eunye Koh, and Handong Zhao. 2023. SpotLight: Visual Insight Recommendation. In *Companion Proceedings of the ACM Web Conference 2023, WWW 2023, Austin, TX, USA, 30 April 2023 - 4 May 2023*, Ying Ding, Jie Tang, Juan F. Sequeda, Lora Aroyo, Carlos Castillo, and Geert-Jan Houben (Eds.). ACM, 19–23. doi:10.1145/3543873.3587302
- [31] Shiyi He, Alexandra Meliou, and Anna Fariha. 2025. ChARLES: Change-Aware Recovery of Latent Evolution Semantics in Relational Data. In *Companion of the 2025 International Conference on Management of Data, SIGMOD/PODS 2025, Berlin, Germany, June 22-27, 2025*, Volker Markl, Joseph M. Hellerstein, and Azza Abouzied (Eds.). ACM, 119–122. doi:10.1145/3722212.3725089
- [32] Wassily Hoeffding. 1994. *Probability Inequalities for sums of Bounded Random Variables*. Springer New York, New York, NY, 409–426. doi:10.1007/978-1-4612-0865-5_26
- [33] Apple Inc. 2025. Numbers - Apple (IN). <https://www.apple.com/in/numbers/> Accessed: 2025-01-21.
- [34] Mahmood Jasim, Christopher Collins, Ali Sarvghad, and Narges Mahyar. 2022. Supporting Serendipitous Discovery and Balanced Analysis of Online Product Reviews with Interaction-Driven Metrics and Bias-Mitigating Suggestions. In *CHI '22: CHI Conference on Human Factors in Computing Systems, New Orleans, LA, USA, 29 April 2022 - 5 May 2022*, Simone D. J. Barbosa, Cliff Lampe, Caroline Appert, David A. Shamma, Steven Mark Drucker, Julie R. Williamson, and Koji Yatani (Eds.). ACM, 9:1–9:24. doi:10.1145/3491102.3517649
- [35] Manas Joglekar, Hector Garcia-Molina, and Aditya G. Parameswaran. 2019. Interactive Data Exploration with Smart Drill-Down. *IEEE Trans. Knowl. Data Eng.* 31, 1 (2019), 46–60. doi:10.1109/TKDE.2017.2685998
- [36] Changhee Joo, Xiaojun Lin, Jiho Ryu, and Ness B. Shroff. 2015. Distributed greedy approximation to maximum weighted independent set for scheduling with fading channels. *IEEE/ACM Transactions on Networking* 24, 3 (2015), 1476–1488.
- [37] Ralph Kimball. 2013. *The data warehouse toolkit: the definitive guide to dimensional modeling* (third edition. ed.). John Wiley & Sons, Incorporated.
- [38] Doris Jung Lin Lee, Dixin Tang, Kunal Agarwal, Thyne Boonmark, Caitlyn Chen, Jake Kang, Ujjaini Mukhopadhyay, Jerry Song, Micah Yong, Marti A. Hearst, and Aditya G. Parameswaran. 2021. Lux: Always-on Visualization Recommendations for Exploratory Dataframe Workflows. *Proc. VLDB Endow.* 15, 3 (2021), 727–738. doi:10.14778/3494124.3494151
- [39] Liangda Li, Hongbo Deng, Anlei Dong, Yi Chang, Ricardo Baeza-Yates, and Hongyuan Zha. 2017. Exploring Query Auto-Completion and Click Logs for Contextual-Aware Web Search and Query Suggestion. In *Proceedings of the 26th International Conference on World Wide Web, WWW 2017, Perth, Australia, April 3-7, 2017*, Rick Barrett, Rick Cummings, Eugene Agichtein, and Evgeniy Gabrilovich (Eds.). ACM, 539–548. doi:10.1145/3038912.3052593
- [40] Qian Liu, Bei Chen, Jiaqi Guo, Morteza Ziyadi, Zeqi Lin, Weizhu Chen, and Jian-Guang Lou. 2022. TAPEX: Table Pre-training via Learning a Neural SQL Executor. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net. <https://openreview.net/forum?id=O50443AsCP>
- [41] Meta Llama. [n.d.]. Llama-3.1-8B-Instruct Model Card. Hugging Face Model Card. <https://huggingface.co/meta-llama/Llama-3.1-8B-Instruct> Accessed 16 July 2025.

- [42] Meta. 2025. Llama3-Instruct. <https://huggingface.co/meta-llama/Meta-Llama-3-8B-Instruct> Accessed: 2025-02-25.
- [43] Microsoft. [n. d.]. Tapex-large Model Card. Hugging Face Model Card. <https://huggingface.co/microsoft/tapex-large> Accessed 16 July 2025.
- [44] Microsoft. 2025. Microsoft Excel - Spreadsheet Software. <https://www.microsoft.com/en-us/microsoft-365/excel> Accessed: 2025-01-21.
- [45] Microsoft Corporation. 2025. Microsoft Power BI. <https://www.microsoft.com/en-us/power-platform/products/power-bi> Accessed: 2025-02-16.
- [46] George A Miller. 1956. The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological review* 63, 2 (1956), 81.
- [47] MOOC.org. n.d. How Important is Excel in Business? <https://www.mooc.org/blog/how-important-is-excel-in-business> Accessed: 2025-01-21.
- [48] Zafeiria Moumoulidou, Andrew McGregor, and Alexandra Meliou. 2021. Diverse Data Selection under Fairness Constraints. In *24th International Conference on Database Theory, ICDT 2021, March 23-26, 2021, Nicosia, Cyprus (LIPIcs, Vol. 186)*, Ke Yi and Zhewei Wei (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 13:1–13:25. doi:10.4230/LIPICS.ICDT.2021.13
- [49] Sk Md Abu Nayeem and Madhumangal Pal. 2007. Genetic algorithmic approach to find the maximum weight independent set of a graph. *Journal of Applied Mathematics and Computing* 25 (2007), 217–229.
- [50] OpenAI. 2025. ChatGPT. <https://chat.openai.com/> Accessed: 2025-02-14.
- [51] Zening Qu and Jessica Hullman. 2018. Keeping Multiple Views Consistent: Constraints, Validations, and Exceptions in Visualization Authoring. *IEEE Trans. Vis. Comput. Graph.* 24, 1 (2018), 468–477. doi:10.1109/TVCG.2017.2744198
- [52] Filip Radlinski and Susan T. Dumais. 2006. Improving personalized web search using result diversification. In *SIGIR 2006: Proceedings of the 29th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval, Seattle, Washington, USA, August 6-11, 2006*, Efthimis N. Efthimiadis, Susan T. Dumais, David Hawking, and Kalervo Järvelin (Eds.). ACM, 691–692. doi:10.1145/1148170.1148320
- [53] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2020. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *J. Mach. Learn. Res.* 21 (2020), 140:1–140:67. <https://jmlr.org/papers/v21/20-074.html>
- [54] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, EMNLP-IJCNLP 2019, Hong Kong, China, November 3-7, 2019*, Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan (Eds.). Association for Computational Linguistics, 3980–3990. doi:10.18653/V1/D19-1410
- [55] sahilnbajaj. [n. d.]. Marketing Campaigns Data Set. Kaggle Dataset. <https://www.kaggle.com/datasets/sahilnbajaj/marketing-campaigns-data-set> Accessed 16 July 2025.
- [56] Shuichi Sakai, Mitsunori Togasaki, and Koichi Yamazaki. 2003. A note on greedy algorithms for the maximum weighted independent set problem. *Discrete applied mathematics* 126, 2-3 (2003), 313–322.
- [57] Salesforce. 2024. Tableau. <https://www.microsoft.com/en-us/power-platform/products/power-bi> <https://www.tableau.com>
- [58] Sunita Sarawagi. 2000. User-Adaptive Exploration of Multidimensional Data. In *Vldb 2000, Proceedings of 26th International Conference on Very Large Data Bases, September 10-14, 2000, Cairo, Egypt*, Amr El Abbadi, Michael L. Brodie, Sharma Chakravathy, Umeshwar Dayal, Nabil Kamel, Gunter Schlageter, and Kyu-Young Whang (Eds.). Morgan Kaufmann, 307–316. <http://www.vldb.org/conf/2000/P307.pdf>
- [59] Sunita Sarawagi. 2001. User-cognizant multidimensional analysis. *Vldb J.* 10, 2-3 (2001), 224–239. doi:10.1007/s007780100046
- [60] Aofeng Su, Aowen Wang, Chao Ye, Chen Zhou, Ga Zhang, Guangcheng Zhu, Haobo Wang, Haokai Xu, Hao Chen, Haoze Li, Haoxuan Lan, Jiaming Tian, Jing Yuan, Junbo Zhao, Junlin Zhou, Kaizhe Shou, Liangyu Zha, Lin Long, Liyao Li, Pengzuo Wu, Qi Zhang, Qingyi Huang, Saisai Yang, Tao Zhang, Wentao Ye, Wufang Zhu, Xiaomeng Hu, Xijun Gu, Xinjie Sun, Xiang Li, Yuhang Yang, and Zhiqing Xiao. 2024. TableGPT2: A Large Multimodal Model with Tabular Data Integration. arXiv:2411.02059 [cs.LG] <https://arxiv.org/abs/2411.02059>
- [61] Aaqib Syed, Phillip Huang Guo, and Vijaykaarti Sundarapandian. 2023. Prune and Tune: Improving Efficient Pruning Techniques for Massive Language Models. In *The First Tiny Papers Track at ICLR 2023, Tiny Papers @ ICLR 2023, Kigali, Rwanda, May 5, 2023*, Krystal Maughan, Rosanne Liu, and Thomas F. Burns (Eds.). OpenReview.net. <https://openreview.net/forum?id=cKlgcx7nSZ>
- [62] Manasi Vartak, Sajjadur Rahman, Samuel Madden, Aditya G. Parameswaran, and Neoklis Polyzotis. 2015. SEEDB: Efficient Data-Driven Visualization Recommendations to Support Visual Analytics. *Proc. VLDB Endow.* 8, 13 (2015), 2182–2193. doi:10.14778/2831360.2831371
- [63] Zhongwei Wan, Xin Wang, Che Liu, Samiul Alam, Yu Zheng, Jiachen Liu, Zhongnan Qu, Shen Yan, Yi Zhu, Quanlu Zhang, Mosharaf Chowdhury, and Mi Zhang. 2024. Efficient Large Language Models: A Survey. *Transactions on Machine Learning Research* (2024). <https://openreview.net/forum?id=bsCCJHbO8A> Survey Certification.
- [64] Zhewei Wei and Ke Yi. 2011. Beyond simple aggregates: indexing for summary queries. In *Proceedings of the 30th ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems, PODS 2011, June 12-16, 2011, Athens, Greece*, Maurizio Lenzerini and Thomas Schwentick (Eds.). PODS, 117–128. doi:10.1145/1989284.1989299
- [65] Junjie Xing, Xinyu Wang, and H. V. Jagadish. 2024. Data-Driven Insight Synthesis for Multi-Dimensional Data. *Proc. VLDB Endow.* 17, 5 (2024), 1007–1019. <https://www.vldb.org/pvldb/vol17/p1007-xing.pdf>
- [66] Cong Yan and Yeye He. 2020. Auto-Suggest: Learning-to-Recommend Data Preparation Steps Using Data Science Notebooks. In *Proceedings of the 2020 International Conference on Management of Data, SIGMOD Conference 2020, online conference [Portland, OR, USA], June 14-19, 2020*, David Maier, Rachel Pottinger, AnHai Doan, Wang-Chiew Tan, Abdussalam Alawini, and Hung Q. Ngo (Eds.). ACM, 1539–1554. doi:10.1145/3318464.3389738
- [67] Cong Yan, Yin Lin, and Yeye He. 2023. Predicate Pushdown for Data Science Pipelines. *Proc. ACM Manag. Data* 1, 2 (2023), 136:1–136:28. doi:10.1145/3589281
- [68] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, Zilin Zhang, and Dragomir R. Radev. 2018. Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing, Brussels, Belgium, October 31 - November 4, 2018*, Ellen Riloff, David Chiang, Julia Hockenmaier, and Jun'ichi Tsujii (Eds.). Association for Computational Linguistics, 3911–3921. doi:10.18653/V1/D18-1425
- [69] Ye Zhang, Stephen Roller, and Byron C. Wallace. 2016. MGNC-CNN: A Simple Approach to Exploiting Multiple Word Embeddings for Sentence Classification. In *NAACL HLT 2016, The 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, San Diego California, USA, June 12-17, 2016*, Kevin Knight, Ani Nenkova, and Owen Rambow (Eds.). The Association for Computational Linguistics, 1522–1527. doi:10.18653/V1/N16-1178