

UTOPIA: Automatic Pivot Table Assistant

Whanhee Cho
University of Utah
Salt Lake City, Utah, USA
whanhee@cs.utah.edu

Anna Fariha
University of Utah
Salt Lake City, Utah, USA
afariha@cs.utah.edu

ABSTRACT

Data summarization is required to comprehend large datasets, and aggregations are effective ways to summarize data. A *pivot table* is a mechanism to aggregate numerical attributes grouped by categorical attributes and spreadsheet pivot tables are particularly suitable for novices, where they can rearrange, group, and aggregate data using intuitive interfaces. However, real-world spreadsheet data is often disorganized, with attributes having *multiple values* and *synonymous variants*. For instance, in the IMDb data, multiple genres are stored as a comma-separated value (“Action, Comedy, Drama”) and “Science Fiction” can be represented in various ways: “Sci-Fi”, “scifi”, “Technological Fiction”, etc. Such data issues pose barriers for novices while constructing pivot tables, and result in noisy and incomprehensible summarization. Parsing multi-valued attributes forces users to resort to external tools (Power Query) or methods (Python), requiring additional expertise; and consolidating synonymous variants demand manual effort, which is a tedious task.

We introduce UTOPIA, an aUTOmatic PIvot table Assistant that extends the functionality of spreadsheet pivot tables, overcoming data issues such as multi-valued attributes and synonymous variants. UTOPIA helps construct pivot tables without requiring additional expertise by (1) automatically detecting multi-valued attributes and organizing the values, achieving *implicit data normalization*, and (2) leveraging SimCSE embeddings and K-Means clustering to consolidate synonymous variants, enabling *semantic aggregation*. We will demonstrate how UTOPIA enables effective pivot table construction, relieving technical novices from tedious data preprocessing while allowing them to remain in their familiar spreadsheet environment, without requiring external tools or additional expertise.

1 INTRODUCTION

The recent growth of data volume and data utilization in machine learning demand effective mechanisms for data analysis and summarization. Aggregation is one of the most effective summarization methods, which involves many operations, such as “sum” over numerical attributes or “count” over categorical attributes. Mechanisms for creating aggregations range from SQL, which requires expert proficiency, to spreadsheet pivot tables, accessible by novices. SQL is suitable for experts, where (1) data must be clean, normalized, and stored in a relational format and (2) a query must be formed using aggregation operators (e.g., SUM, AVG, CNT). Alternatively, novices can utilize spreadsheet pivot table (Table 1(c)), which offers an intuitive and easy interface to group data by certain attributes (specified as “row”) and apply aggregates to other attributes of interest (specified by “column”).

Spreadsheets such as Microsoft Excel¹ and Google Sheets² offer pivot table functionality, where users can use non-relational

or “flat” data, and perform relatively straightforward data analysis. Despite being designed for novices, spreadsheet pivot tables bring adversities when the data is disorganized and messy, preventing the users from obtaining the desired results directly. Disorganized data requires careful, often manual, pre-processing, which demands additional technical expertise—such as programming in Python or Power Query³—and time-consuming manual effort such as identifying all synonyms and replacing them with a canonical value. Technical novices comprise approximately 11% of office workers [1], who lack programming expertise (Python, PowerQuery, Regex, etc.). While business intelligence models such as Tableau⁴ and Microsoft Power BI target such novices, they still require the user to learn Power Query or Regex to perform nontrivial data preprocessing, and offer no solution to automatically consolidate synonymous variants.

EXAMPLE 1. *Patel, a technical novice, is working with the IMDb⁵ data about 1000 most popular movies (Table 1(a)). Patel wants to know which film genres yield significant gross. She decides to use Microsoft Excel pivot table, grouping by Genre and summing over Gross, and expects to obtain Table 1(c). To her disappointment, she gets Table 1(b). The reason is that Genre contains multiple values for some movies, such as “action, crime, drama”, and Excel incorrectly assumes that this entire comma-separated list is the value for genre. To get her desired result, Patel must parse the multi-valued attribute Genre, which requires using an advanced Excel functionality or writing a Python script.*

However, Patel’s struggle continues even after parsing these comma-separated lists. She realizes that some movies have only one genre, whereas some have up to 10 different genres. She must decide how to store such variable-length values against a single attribute Genre. She can derive 10 different attributes labeled by Genre1, Genre2, ..., Genre10, but this will result into two issues. First, many cells will be empty as not all movies have 10 genres, resulting in poor data representation. Second, she will still struggle to obtain her desired pivot table (Table 1(c)) as the target attribute is now split across 10 different attributes. If she generates one pivot table for each of the newly generated attributes, it will result in 10 different pivot tables, which is not what Patel wants. While Power Query may offer a solution, Patel hesitates to leave the spreadsheet environment, which she is very comfortable with.

EXAMPLE 2. *Patel somehow parsed the multi-valued attribute Genre using an external tool (Microsoft Power Query) and is now working on a different movie dataset (Table 2(a)). She proceeds to generate a pivot table over this new dataset and obtains Table 2(b). Patel struggles to interpret the results as she expects “action” to be the top-gross genre, but the result indicates “biography” to have a*

¹Microsoft www.microsoft.com/en-us/microsoft-365/excel

²Google Sheets: www.google.com/sheets/about/

³Microsoft Power Query powerquery.microsoft.com/en-us/

⁴Tableau: www.tableau.com/

⁵IMDB: www.kaggle.com/datasets/PromptCloudHQ/imdb-data

Title	Genre	Gross	Genre	Sum of Gross	Genre	Sum of Gross
Joker	drama	28 M	action, crime, drama	535 M	action	1322 M
2001: A Space Odyssey	action, crime, drama	535 M	action, sci-fi	464 M	drama	697 M
Queen	action, sci-fi	171 M	action, adventure	323 M	crime	535 M
The Prestige	biography, drama	97 M	biography, drama	97 M	sci-fi	464 M
The Departed	action, sci-fi	293 M	drama	65 M	adventure	323 M
The Usual Suspects	drama	37 M			biography	97 M
Back to the Future	action, adventure	323 M	(b)		(c)	

Table 1: (a) An sample from the IMDb dataset. **Title** indicates title of a movie, **Genre** indicates the movie’s genres, separated by commas, and **Gross** indicates the corresponding profit. (b) Since **Genre** contains multiple values, if the values are not parsed before pivot table construction, such ill-formed pivot table is produced. (c) The desired pivot table.

Title	Genre	Gross	Genre	Sum of Gross	Genre	Sum of Gross
The Shawshank Redemption	prison drama	28 M	biography	630 M	action	1029 M
The Dark Knight	superhero action, crime, epic drama	535 M	epic drama	535 M	biography	630 M
The Matrix	action, epic sci-fi	171 M	crime	535 M	drama	600 M
Schindler’s List	biography	630 M	superhero action	535 M	crime	535 M
Inception	team action, space opera sci-fi	93 M	epic action	323 M	adventure	323 M
Fight Club	drama	37 M	space adventure	323 M	sci-fi	264 M
Star Wars	epic action, space adventure	323 M	action	171 M	(c)	
(a)				
			(b)			

Table 2: (a) An example dataset from IMDb. (b) Interpreting this pivot table is challenging due to the presence of synonymous variants in **Genre**, such as misspellings and sub-genres. (c) The expected pivot table.

larger gross than “action”. With a closer look, she realizes that the data has different representations for the same genres: “action” has synonymous variants such as “superhero action”, “team action”, and “epic action”. Patel wants to generate Table 2(c), but for that, she must replace all variants of similar semantics with a canonical value, which demands domain expertise and is a tedious process.

Patel’s examples show that generating a pivot table over disorganized data requires tedious data preprocessing. While some tools offer mechanisms for parsing multi-valued attributes, they often force users to depart from the spreadsheet environment and resort to external tools requiring specialized skills. For parsing multiple values, Microsoft Excel requires Power Query, Tableau requires use of regex grammar, and user forums recommend programming using spreadsheet functions or Python. While valid, all these mechanisms require programming expertise, which novices lack. Even for a simple parsing problem like the one shown in Example 1, in Python, one needs to write code involving six steps: (1) loading the CSV file in a DataFrame object, (2) splitting text, (3) iteratively creating new attribute names based on the maximum number of fragments obtained from splitting, (4) concatenating the new attributes to the original DataFrame, (5) removing the original attribute to avoid redundancy, and (6) storing the preprocessed data back to the CSV format. In contrast, relational databases expect normalization: a “genre” relation can be introduced to contain pairs of movie and a single genre in each row, which requires understanding normalization and restructuring flat data format to relational format that is not novice-friendly.

Handling synonymous variants requires domain knowledge and entails tedious work as users must explore and modify variants manually. For instance, in Example 2, if there are on average 25 variants for each of the 20 genres, one will have to manually issue 500 find-replace operations! Moreover, this will fail for typos, case mismatches, and formatting issues. Synonymous variants mandate

that users thoroughly explore the entire data to determine which values have equivalent semantic meanings. Another critical issue is that such an operation is *irreversible*, the user will lose the information about variants once they consolidate synonymous variants explicitly.

Therefore, an ideal pivot table functionality should be able to handle disorganized data gracefully and have three key features: (1) it should be able to extract values from multi-valued attributes, (2) it should be able to perform *semantic aggregation*, by grouping synonymous variants *implicitly*, without altering the data, and (3) it should let the user stay in their preferred spreadsheet environment.

We introduce UTOPIA—an **uTOMatic PI**ivot table **A**ssistant—to extend the functionality of spreadsheet pivot tables, overcoming data issues like multi-valued attributes and synonymous variants. UTOPIA automatically parses and organizes values for multi-valued attributes, providing *implicit* data normalization. Additionally, UTOPIA enables *semantic aggregation* using sentence embeddings and K-Means clustering to consolidate synonymous variants. Finally, UTOPIA produces a *dynamic* and *interactive* pivot table with options for expanding and collapsing data values to display synonymous variants.

Related work. Microsoft Excel FlashFill and Google Sheets Smart-Fill can split multi-valued attributes but require manual input of examples with uniform delimiters and knowledge of the resulting fragments, demanding significant user effort. Unlike UTOPIA, these tools cannot handle synonymous variants. Auto-Tables [4] normalizes multi-valued attribute by keeping only the first value and disregarding the rest, addressing only the structural issue of the data, but neglecting its content. Automatic data analysis tools [3, 5, 9] assist in data summarization, but they focus only on summarization, while UTOPIA aims to reduce the barrier of data preprocessing for technical novices, specifically for the task of pivot table generation in a spreadsheet environment.

In our demonstration, participants will observe how UTOPIA effectively generates the desired pivot table, bypassing data issues such as multi-valued attributes and synonymous variants, and without requiring any additional effort from the user. We proceed to provide an overview of our system in Section 2 and then a detailed walkthrough of our demonstration scenario, based on Example 2, in Section 3.

2 SYSTEM OVERVIEW

UTOPIA consists of three key components: (1) Multi-Valued Attribute Handler, which automatically identifies and parses multiple values within an attribute, (2) Synonymous Variants Handler, which identifies and semantically aggregates variants, and (3) Data Organizer, which generates an interactive pivot table.

Multi-Valued Attribute Handler. This component is responsible for detecting and parsing multi-valued attributes, thereby achieving data normalization. UTOPIA can identify multiple values, even if explicit delimiters are absent. For instance, for the value “superhero actioncrimeepic drama”, UTOPIA recognizes the presence of multiple values and parses them by transforming them into a set: {“superhero action”, “crime”, “epic drama”}. To accomplish this, UTOPIA leverages a data extraction algorithm [7] that enables the extraction of multiple values even in the absence of explicit delimiters.

Synonymous Variants Handler. Next, UTOPIA searches for synonymous variants within each relevant attribute. To this end, UTOPIA computes the ratio r between the number of unique and total values. If r exceeds a certain threshold $\tau_{lb} \in [0, 1]$, UTOPIA determines that synonymous variants *may* exist. E.g., if there are semantically related values such as “Sci-Fi”, “Science Fiction”, and “Cyberpunk”, then a large number of unique values will exist, leading to a large r . However, a very high value for r indicates that observing unique values is natural for that attribute. For instance, people’s first names may have many unique values, but they are not synonymous variants. Therefore, when r does not exceed the threshold $\tau_{ub} \in [0, 1]$ (while exceeding τ_{lb}), UTOPIA assumes presence of synonymous variants. While the user is free to tune the parameters τ_{lb} and τ_{ub} , we found the values $\tau_{lb} = 0.4$ and $\tau_{ub} = 0.8$ to work well in practice.

To find values that share similar or identical semantic meanings, UTOPIA employs word embeddings to compute the similarity between values in the vector space. For instance, in Table 2, “action” and “superhero action” share the same semantic meaning, resulting in a small distance between them in the embedding space. Since traditional word embeddings [6] struggle to comprehend long phrases such as “space opera sci-fi”, UTOPIA opts for state-of-the-art pre-trained sentence embedding model, SimCSE [2], which obtains superior performance over Sentence-BERT [8].

Next, UTOPIA consolidates semantically similar values using K-Means clustering. It suggests the value of k based on the best *silhouette score*, a specialized metric to measure cluster quality that is convex-shaped. However, k is a customizable parameter and the user can tune it according to their preferences. For instance, if fine-grained grouping is desired, k should be set to a large value, and for a more generalized view of the data, k should be set to a smaller value. Note that traditional relational databases do not offer such flexibility due to strict schema requirements.

Data Organizer. Finally, UTOPIA displays a pivot table with parsed data and semantically aggregated variants. UTOPIA represents parsed

values as row or column labels in the pivot table. UTOPIA displays a representative value (e.g., “action” is chosen as representative for “action”, “superhero action”, “team action”, etc.) in the presence of synonymous variants. The representative value is the one with the closest embedding to the average embeddings over all variants.

To ensure that the original data integrity is preserved without any loss of information, UTOPIA stores the parsed data in JSON format, a natural choice over tabular format for storing multi-valued attributes. This also avoids repetitive parsing computation for subsequent operations. Moreover, UTOPIA is robust to data updates: instead of recomputing clusters for minor data changes, it assigns the new data to the most similar cluster. Note that our focus is on *usability* enhancement and our target platform is spreadsheets, which doesn’t support big data, indicating no scalability challenge is involved.

Preliminary results: Over the IMDb dataset and with 27 expected genres, we found UTOPIA to attain a cluster purity score—which indicates how much each cluster contains semantically similar values—of 0.86. We also tried ChatGPT 3.5 using the prompt “Group the following words into semantically related groups. Don’t change or omit words. Create k groups.” for different values of k . However, ChatGPT shows an interesting (but undesirable) behavior when k is smaller than ideal. With $k = 10$, ChatGPT forms the following groups: {“Action and Adventure”, “Animation and Comedy”, “Crime and Documentary”, ...}. While ChatGPT succeeds in including all sub-genres of “Action” and “Adventure” into “Action and Adventure”, it incorrectly merges groups based on their lexicographic similarity: “Action” is alphabetically close to “Adventure” but not semantically. Ideally, one would prefer “Action” to be merged with “Thriller” or “Crime” over “Adventure”. In contrast UTOPIA provides semantically meaningful groups even when fewer clusters are requested by the user, usually for enhanced interpretability.

3 DEMONSTRATION SCENARIO

We will demonstrate UTOPIA over datasets from various domains, such as recipe data listing multiple ingredients⁶, IMDb data including various sub-genres, and university course survey over various departments⁷. Below we provide a demonstration scenario over part of the IMDb dataset, which contains metadata for the top 1000 successful movies, with eight attributes (movie title, year, genre, etc.). We randomly introduced some misspellings and augmented this data with sub-genres. We will guide users through eleven steps (annotated in Figure 1) impersonating Patel, who is interested in generating a pivot table to find top-grossing movie genres for each year.

Step A (Uploading data) The user first uploads a (disorganized) data containing multi-valued attributes and synonymous variants. In this case, the IMDb dataset for top 1000 successful movies.

Step B (Viewing data overview) Next, the user previews the data. UTOPIA displays the first three rows. The user can horizontally or vertically scroll to see more data.

Step C (Selecting attributes) Then the user selects the attributes they want to focus on for the pivot table generation. In our scenario, the user chooses Year, Genre, and Gross.

Step D (Choosing positions) The user can assign a selected attribute to one of the pivot table positions, “Column”, “Row”, or

⁶Recipe dataset: www.github.com/majumderb/recipe-personalization

⁷www.kaggle.com/datasets/sank3t/university-student-survey

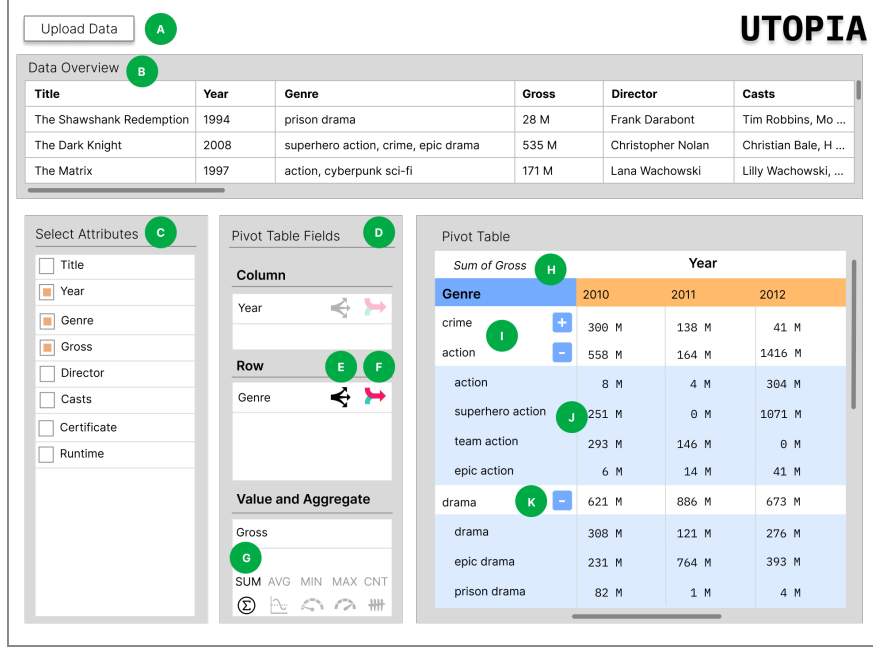


Figure 1: The UTOPIA demo: (A) upload disorganized data, (B) view data overview, (C) select attributes to analyze, (D) choose a position for the selected attribute for the pivot table, (E) select multi-valued attribute handler, (F) select synonymous variants handler, (G) choose aggregation method, (H) view aggregated attribute, (I) view the parsed data of multi-valued attributes, (J) view semantically aggregated values from synonymous variants, (K) extend or collapse semantically aggregated variants.

“Value”, by dragging it to the desired position. When an attribute moves to “Row” or “Column”, then the values of that attribute will be assigned as row or column labels in the pivot table, respectively. If the user selects an attribute to “Value”, it will be aggregated. In our guided scenario, the user chooses *Year* for column, *Genre* for row, and *Gross* for value.

Step (E) (Enabling multi-valued attribute handler) UTOPIA displays an icon, next to each row/column attribute, representing multi-valued attribute handler. This icon is gray (disabled) if the attribute does not contain multiple values. In our guided scenario, *Year* has a disabled icon, while *Genre* has an enabled icon. The user chooses to keep it enabled and proceeds.

Step (F) (Enabling synonymous variants handler) Next to the multi-valued attribute handler icon, UTOPIA displays another icon, representing synonymous variants handler. This icon is gray (disabled) if the attribute does not contain synonymous variants. Otherwise, it is enabled. The user can choose to disable it if they wish. In our guided scenario, *Year* has a disabled icon, while *Genre* has an enabled icon. The user chooses to keep it enabled and proceeds. Moreover, by right-clicking this icon, the user can specify additional system parameters τ_{lb} , τ_{ub} , and k (not shown in the figure).

Step (G) (Choosing an aggregation method) The user selects SUM as the aggregation method over *Gross*.

Step (H) (Viewing the pivot table) UTOPIA now produces a pivot table in the bottom right panel. All values for *Gross* for the *Genre* “action” are aggregated together: showing the sum of gross as 558 M for 2010.

Step (I) (Viewing parsed data from multi-valued attribute) UTOPIA automatically parses the values within *Genre* and organizes the resulting values as row labels.

Step (J) (Viewing semantic aggregation) UTOPIA semantically aggregates the synonymous variants and shows the representative value on top. E.g., “action” has four synonymous variants.

Step (K) (Expanding or collapsing variants) The user can expand (collapse) the row labels to show (hide) the synonymous variants. E.g., expanding “action” reveals four synonymous variants.

After the guided demonstration, participants may use UTOPIA to explore their own datasets. The key takeaway is the convenience UTOPIA provides for creating pivot tables over disorganized data, without requiring any additional skill from the user. In summary, UTOPIA targets technical novices who seek to avoid technically challenging and time-consuming data preprocessing while generating pivot tables in a spreadsheet environment.

REFERENCES

- [1] 2024. Excel Facts & Statistics: New Research 2024. www.acuitytraining.co.uk/news-tips/new-excel-facts-statistics-2022/.
- [2] T Gao, X Yao, and D Chen. 2021. SimCSE: Simple Contrastive Learning of Sentence Embeddings. In *EMNLP*. 6894–6910.
- [3] R J L John, D Bacon, J Chen, U Ramesh, J Li, D Das, R V Claus, A Kendall, and J M Patel. 2023. DataChat: An Intuitive and Collaborative Data Analytics Platform. In *SIGMOD*. 203–215.
- [4] P Li, Y He, C Yan, Y Wang, and S Chaudhuri. 2023. Auto-Tables: Synthesizing Multi-Step Transformations to Relationalize Tables without Using Examples. *PVLDB* 16 (2023), 3391–3403.
- [5] P Ma, R Ding, S Wang, SHan, and D Zhang. 2023. InsightPilot: An LLM-Empowered Automated Data Exploration System. In *EMNLP*. 346–352.
- [6] T Mikolov, K Chen, G Corrado, and J Dean. 2013. Efficient Estimation of Word Representations in Vector Space. In *ICLR*.
- [7] M Raza and S Gulwani. 2017. Automated Data Extraction Using Predictive Program Synthesis. In *AAAI*. 882–890.
- [8] N Reimers and I Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *EMNLP*. 3980–3990.
- [9] A Wu, Y Wang, M Zhou, X He, H Zhang, H Qu, and D Zhang. 2022. MultiVision: Designing Analytical Dashboards with Deep Learning Based Recommendation. *EEE Trans Vis Comput Graph* 28, 1 (2022), 162–172.